

# SCS\_3253\_030\_ML\_Project\_Final-2

July 22, 2019

## 1 PROBLEM DEFINITION

**DATASET:** This dataset comprises of data from Bike Sharing Company related to Bike usage over the year 2011 & 2012 in Capital bikeshare system in Washington, DC .

**PROBLEM DEFINITION:** This is a supervised learning exercise and regression or classification models can be used to solve the problem.

Regression: Predict demand for Bike in particular time of the year.

Both hour.csv and day.csv have the following fields, except hr which is not available in day.csv

- instant: record index
- dteday : date
- season : season (1:summer, 2:summer, 3:fall, 4:winter)
- yr : year (0: 2011, 1:2012)
- mnth : month ( 1 to 12)
- hr : hour (0 to 23)
- holiday : weather day is holiday or not (extracted from [Web Link])
- weekday : day of the week
- workingday : if day is neither weekend nor holiday is 1, otherwise is 0.
- weathersit :
- 1: Clear, Few clouds, Partly cloudy, Partly cloudy
- 2: Mist + Cloudy, Mist + Broken clouds, Mist + Few clouds, Mist
- 3: Light Snow, Light Rain + Thunderstorm + Scattered clouds, Light Rain + Scattered clouds
- 4: Heavy Rain + Ice Pallets + Thunderstorm + Mist, Snow + Fog
- temp : Normalized temperature in Celsius. The values are derived via  $(t-t_{min})/(t_{max}-t_{min})$ ,  $t_{min}=-8$ ,  $t_{max}=+39$  (only in hourly scale)
- atemp: Normalized feeling temperature in Celsius. The values are derived via  $(t-t_{min})/(t_{max}-t_{min})$ ,  $t_{min}=-16$ ,  $t_{max}=+50$  (only in hourly scale)
- hum: Normalized humidity. The values are divided to 100 (max)
- windspeed: Normalized wind speed. The values are divided to 67 (max)
- casual: count of casual users
- registered: count of registered users
- cnt: count of total rental bikes including both casual and registered

## 2 IMPORT FUNCTIONALITY

```
In [0]: import pandas as pd
import numpy as np
```

```

from sklearn.metrics import mean_squared_error
from sklearn.metrics import (r2_score, explained_variance_score, max_error, mean_absolute_
mean_squared_error, mean_squared_log_error, median_absolute_error)
import seaborn as sns
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import learning_curve
from sklearn.model_selection import validation_curve

# To plot pretty figures
%matplotlib inline
import matplotlib
import matplotlib.pyplot as plt
plt.rcParams['axes.labelsize'] = 14
plt.rcParams['xtick.labelsize'] = 12
plt.rcParams['ytick.labelsize'] = 12

# to make this notebook's output stable across runs
np.random.seed(123)

```

### 3 DATA DOWNLOAD AND REVIEW

Data Set Information:

Bike sharing systems are new generation of traditional bike rentals where whole process from membership, rental and return back has become automatic. Through these systems, user is able to easily rent a bike from a particular position and return back at another position. Currently, there are about over 500 bike-sharing programs around the world which is composed of over 500 thousands bicycles. Today, there exists great interest in these systems due to their important role in traffic, environmental and health issues.

Apart from interesting real world applications of bike sharing systems, the characteristics of data being generated by these systems make them attractive for the research. Opposed to other transport services such as bus or subway, the duration of travel, departure and arrival position is explicitly recorded in these systems. This feature turns bike sharing system into a virtual sensor network that can be used for sensing mobility in the city. Hence, it is expected that most of important events in the city could be detected via monitoring these data.

<https://archive.ics.uci.edu/ml/datasets/Bike+Sharing+Dataset>

```

In [2]: #!git clone https://github.com/emanhamed/Houses-dataset
        !git clone https://github.com/asifsundrani/SCS_3253_030-Machine-Learning-Project

```

```

Cloning into 'SCS_3253_030-Machine-Learning-Project'...
remote: Enumerating objects: 7, done.
remote: Counting objects: 100% (7/7), done.
remote: Compressing objects: 100% (6/6), done.
remote: Total 7 (delta 0), reused 0 (delta 0), pack-reused 0
Unpacking objects: 100% (7/7), done.

```

```

In [0]: df_hour=pd.read_csv('SCS_3253_030-Machine-Learning-Project/hour.csv')
        df_day=pd.read_csv('SCS_3253_030-Machine-Learning-Project/day.csv')

```

```
In [4]: df_day.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 731 entries, 0 to 730
Data columns (total 16 columns):
instant      731 non-null int64
dteday       731 non-null object
season       731 non-null int64
yr           731 non-null int64
mnth         731 non-null int64
holiday      731 non-null int64
weekday      731 non-null int64
workingday   731 non-null int64
weathersit    731 non-null int64
temp         731 non-null float64
atemp        731 non-null float64
hum          731 non-null float64
windspeed    731 non-null float64
casual       731 non-null int64
registered   731 non-null int64
cnt          731 non-null int64
dtypes: float64(4), int64(11), object(1)
memory usage: 91.5+ KB
```

Comments: df\_day is the summary of cumulative hourly event of the day and we won't use this data for modelling due to lower number of samples.

```
In [5]: df_hour.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17379 entries, 0 to 17378
Data columns (total 17 columns):
instant      17379 non-null int64
dteday       17379 non-null object
season       17379 non-null int64
yr           17379 non-null int64
mnth         17379 non-null int64
hr           17379 non-null int64
holiday      17379 non-null int64
weekday      17379 non-null int64
workingday   17379 non-null int64
weathersit    17379 non-null int64
temp         17379 non-null float64
atemp        17379 non-null float64
hum          17379 non-null float64
windspeed    17379 non-null float64
casual       17379 non-null int64
registered   17379 non-null int64
```

```

cnt                17379 non-null int64
dtypes: float64(4), int64(12), object(1)
memory usage: 2.3+ MB

```

All the data fields are in numeric except dteday. Also there is no null value in the dataset. There are certain fields which can be eliminated for the purpose for this exercise a)instant: this is just a reference

b)dteday: date which is not relevant for this exercise

c)yr: not relevant for the purpose of the exercise

d)casual & registered: target value set to total value so need to distribute in casual and registered.

These fields are not relevant for the purpose of this exercise.

Target field is going to cnt i.e. total count of bikers i.e. both registered and casual.

```
In [6]: df_hour.describe()
```

```

Out [6]:
         instant      season  ...  registered      cnt
count  17379.0000  17379.000000  ...  17379.000000  17379.000000
mean    8690.0000    2.501640  ...    153.786869   189.463088
std     5017.0295    1.106918  ...    151.357286   181.387599
min         1.0000    1.000000  ...     0.000000    1.000000
25%     4345.5000    2.000000  ...     34.000000   40.000000
50%     8690.0000    3.000000  ...    115.000000  142.000000
75%    13034.5000    3.000000  ...    220.000000  281.000000
max     17379.0000    4.000000  ...    886.000000  977.000000

```

[8 rows x 16 columns]

Categorical Fields: Following fields can be classified as categorical i.e. season, month, hour, weekday & weathersit. The value associated with them are in range of numbers therefore it would be interesting to see how model behaves with and without onehotcode translation.

Numeric Fields: Numeric fields include weather related characteristics and these have been normalized based on each scenario. Fields include temp, atemp, hum & windspeed.

Target field is cnt i.e. total count of riders

```
In [7]: df_hour.head()
```

```

Out [7]:
   instant  dteday  season  yr  ...  windspeed  casual  registered  cnt
0         1  2011-01-01     1   0  ...      0.0        3          13   16
1         2  2011-01-01     1   0  ...      0.0        8          32   40
2         3  2011-01-01     1   0  ...      0.0        5          27   32
3         4  2011-01-01     1   0  ...      0.0        3          10   13
4         5  2011-01-01     1   0  ...      0.0        0           1    1

```

[5 rows x 17 columns]

```
In [8]: df_day.head()
```

```
Out[8]:
```

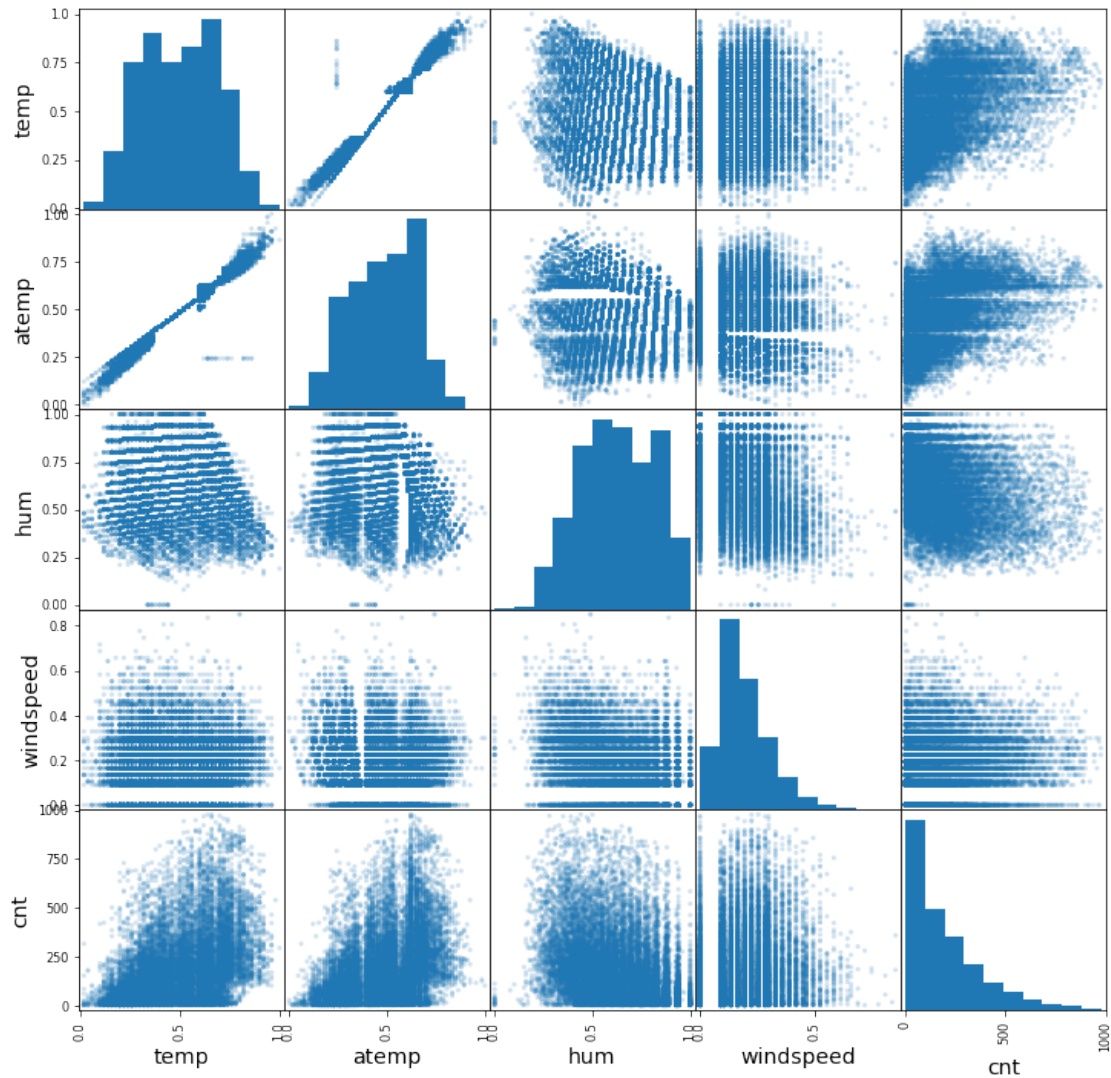
	instant	dteday	season	yr	...	windspeed	casual	registered	cnt
0	1	2011-01-01	1	0	...	0.160446	331	654	985
1	2	2011-01-02	1	0	...	0.248539	131	670	801
2	3	2011-01-03	1	0	...	0.248309	120	1229	1349
3	4	2011-01-04	1	0	...	0.160296	108	1454	1562
4	5	2011-01-05	1	0	...	0.186900	82	1518	1600

[5 rows x 16 columns]

## 4 DATA VISUALIZATION

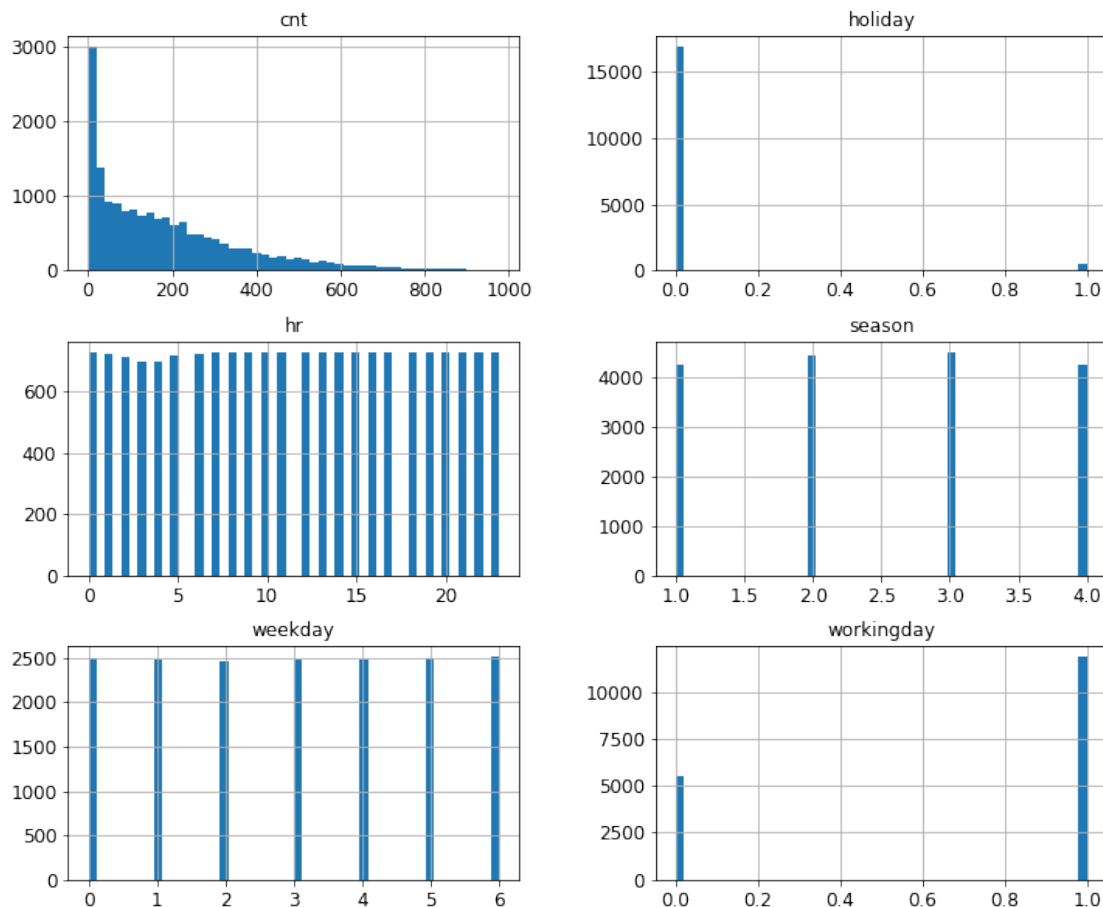
```
In [0]: from pandas.plotting import scatter_matrix
```

```
In [10]: scatter_matrix(df_hour[['temp','atemp','hum','windspeed','cnt']], alpha=0.2, figsize=
save_results_to = ''
plt.savefig(save_results_to + 'scatter matrix', dpi = 300)
```



```
In [0]: df_hour[['season','hr','holiday','weekday','workingday','cnt']].hist(bins=50,figsize=(
```

```
Out[0]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f81bcf5c8d0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f81bccd3358>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x7f81bcc7898>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f81bb48ce10>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x7f81bb43a3c8>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f81bb462940>]],
dtype=object)
```

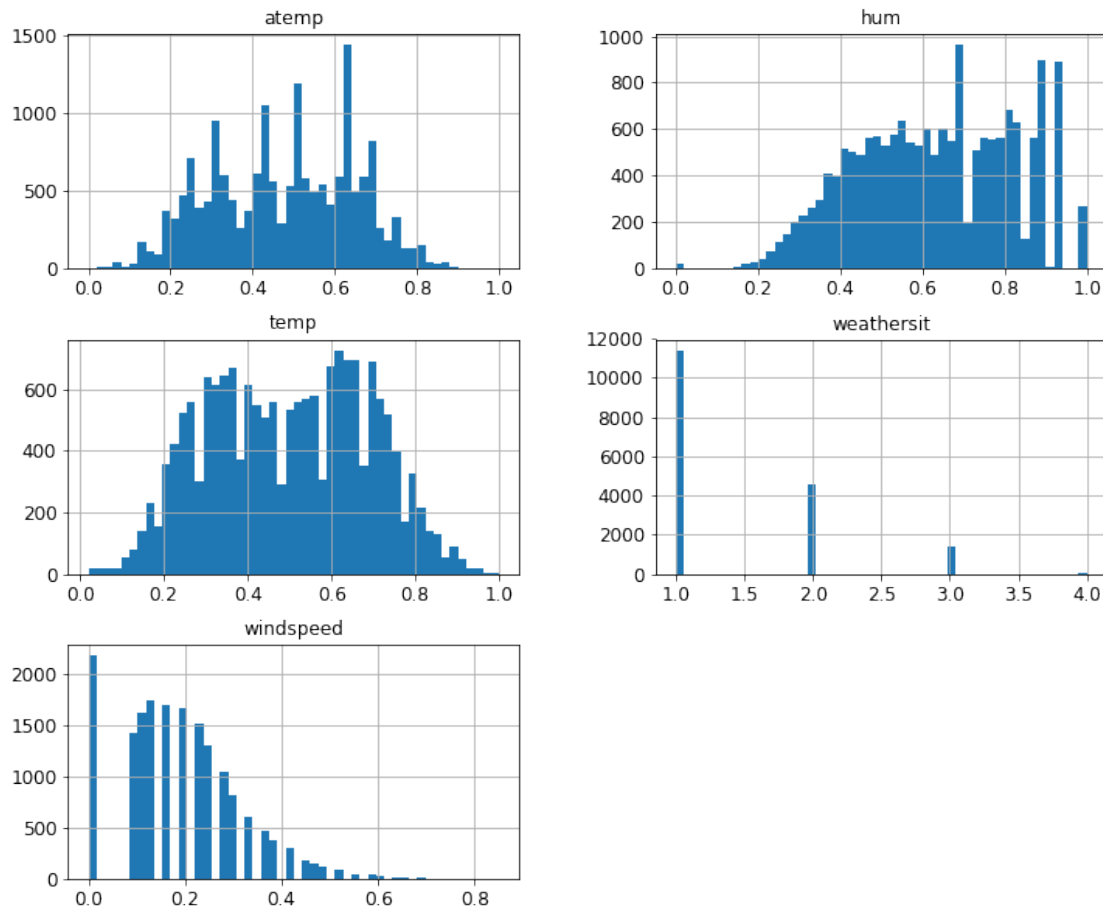


#### Observations:

Interestingly demand for bikes are consistent accross following areas i.e. weekdays, hours, season and working day vs weekend However, lower demand on holidays vs other days (weekdays and weekends)

```
In [0]: df_hour[['temp','atemp','hum','windspeed','weathersit']].hist(bins=50,figsize=(12,10))
```

```
Out[0]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f81bcceb438>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f81bb3aa4a8>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x7f81bb010a20>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f81bafb5f98>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x7f81baf67550>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f81baf8dac8>]],
dtype=object)
```



Observations:

Windspeed: Higher windspeed drives lower demand for bike as shown by the graph.  
Weathersit: Demand for bikes goes down as weather from clear to cloudy and bad weather reduces the demand for the bike.  
temp (Normalized temperature in Celsius) and atemp(Normalized feeling temperature in Celsius) are showing normal distribution curve.  
humidity: interestingly higher humidity encouraged more bikers to bike.

## 5 DATA CLEAN UP

```
In [11]: df_hour.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17379 entries, 0 to 17378
Data columns (total 17 columns):
instant      17379 non-null int64
dteday       17379 non-null object
season       17379 non-null int64
yr           17379 non-null int64
mnth         17379 non-null int64
hr           17379 non-null int64
holiday      17379 non-null int64
weekday      17379 non-null int64
workingday   17379 non-null int64
weathersit    17379 non-null int64
temp         17379 non-null float64
atemp        17379 non-null float64
hum          17379 non-null float64
windspeed    17379 non-null float64
casual       17379 non-null int64
registered   17379 non-null int64
cnt          17379 non-null int64
dtypes: float64(4), int64(12), object(1)
memory usage: 2.3+ MB

```

```
In [12]: df_hour['weekday'].value_counts(sort=False)
```

```

Out[12]: 0    2502
         1    2479
         2    2453
         3    2475
         4    2471
         5    2487
         6    2512
         Name: weekday, dtype: int64

```

Normal distribution of Bikers over the weekdays.

## 6 DATA PREPROCESSING AND PIPELINE

```
In [0]: df_model=df_hour.drop(['instant','dteday','yr','casual','registered'], axis=1)
```

These fields are not relevant related to the problem therefore they are being dropped off from the evaluation.

```
In [0]: df_model.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17379 entries, 0 to 17378

```



```
Data columns (total 12 columns):
season      17379 non-null int64
mnth       17379 non-null int64
hr         17379 non-null int64
holiday    17379 non-null int64
weekday    17379 non-null int64
workingday 17379 non-null int64
weathersit  17379 non-null int64
temp       17379 non-null float64
atemp      17379 non-null float64
hum        17379 non-null float64
windspeed  17379 non-null float64
cnt        17379 non-null int64
dtypes: float64(4), int64(8)
memory usage: 1.6 MB
```

```
In [0]: from sklearn.pipeline import Pipeline
        from sklearn.preprocessing import StandardScaler, OneHotEncoder
        from sklearn.compose import ColumnTransformer
        from sklearn.impute import SimpleImputer

categorical_features = ['season', 'mnth', 'hr', 'weekday', 'weathersit', 'holiday', 'workingday']
categorical_transformer = Pipeline([('onehot', OneHotEncoder(handle_unknown='ignore'))])

numeric_features = ['temp', 'atemp', 'hum', 'windspeed']
numeric_transformer = Pipeline([
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())])

preprocessor = ColumnTransformer([
    ('num', numeric_transformer, numeric_features),
    ('cat', categorical_transformer, categorical_features)
])

#DATA SPLIT - TARGET & TRAIN_TEST_SPLIT

In [15]: X = df_model.drop('cnt', axis=1)
        y = df_model['cnt']
        print(X.shape)
        print(y.shape)

(17379, 11)
(17379,)
```

```
In [16]: from sklearn.model_selection import train_test_split

        X_tr, X_te, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```

print(X_tr.shape,y_train.shape)
print(X_te.shape,y_test.shape)
print('Mean of X_Train:',np.mean(y_train))

```

```

(13903, 11) (13903,)
(3476, 11) (3476,)
Mean of X_Train: 189.41451485290943

```

```

In [17]: X_train=preprocessor.fit_transform(X_tr)
        X_test=preprocessor.transform(X_te)
        print(X_train.shape,y_train.shape)
        print(X_test.shape,y_test.shape)
        print('Mean of X_Train:',np.mean(y_train))

```

```

(13903, 59) (13903,)
(3476, 59) (3476,)
Mean of X_Train: 189.41451485290943

```

```

In [18]: from sklearn.preprocessing import PolynomialFeatures
        d=2
        poly_features = PolynomialFeatures(degree=d, interaction_only=False,include_bias=False)
        X_poly = poly_features.fit_transform(X_train)
        X_poly.shape

```

```

Out[18]: (13903, 1829)

```

```

In [19]: X_poly_test = poly_features.fit_transform(X_test)
        X_poly_test.shape

```

```

Out[19]: (3476, 1829)

```

## 7 SET UP GRID SEARCH AND POLYNOMIAL FEATURES (ADD COMPLEXITY)

```

In [0]: from sklearn.model_selection import GridSearchCV

        def best_parameters(estimator,param_grid,cv):
            gs1=GridSearchCV(estimator,param,cv=cv,verbose=1,n_jobs=-1)
            gs1_results=gs1.fit(X_train,y_train)
            estimator=gs1_results
            print('Best_results:',estimator.best_score_)
            print('Best_parameters:',estimator.best_params_)
            print('Best_estimator:',estimator.best_estimator_)

In [0]: def best_parameters_poly(estimator,param_grid,cv):
        gs1=GridSearchCV(estimator,param,cv=cv,verbose=1,n_jobs=-1)

```

```

gs1_results=gs1.fit(X_poly,y_train)
estimator=gs1_results
print('Best_results:',estimator.best_score_)
print('Best_parameters:',estimator.best_params_)
print('Best_estimator:',estimator.best_estimator_)

```

## 8 SET UP OF METRICS AND EVALUATION :

In [0]: # PERFORMANCE METRICS FOR REGRESSION METHODOLOGY

```

def lin_perform_metrics(target,prediction):
    print('R2_Square: ',r2_score(target,prediction))
    print('Explained_Variance_Score: ',explained_variance_score(target,prediction))

    print('Mean_absolute_error: ',mean_absolute_error(target,prediction))
    print('Median_absolute_error: ',median_absolute_error(target,prediction))
    print('Max_error: ',max_error(target,prediction))

    MSE=mean_squared_error(target,prediction)
    print('MSE: ',MSE)
    print('RMSE: ',np.sqrt(MSE))

def lin_peform_metrics_log(target,prediction):
    print('Mean_squared_log_error: ',mean_squared_log_error(target,prediction))

```

In [0]: # LEARNING CURVE FULL DATASET EVALUATION

```

def plot_learning_curve_full(model,X,y):
    X_train,X_val,y_train,y_val=train_test_split(X,y,test_size=0.2)
    train_errors,val_errors=[],[]
    for m in range(1,len(X_train)):
        model.fit(X_train[:m],y_train[:m])
        y_train_predict=model.predict(X_train[:m])
        y_val_predict=model.predict(X_val)
        train_errors.append(mean_squared_error(y_train[:m],y_train_predict))
        val_errors.append(mean_squared_error(y_val, y_val_predict))
    plt.plot(np.sqrt(train_errors),'r--',linewidth=2,label='Training MSE Score')
    plt.plot(np.sqrt(val_errors),'b-',linewidth=2,label='Validation MSE Score')

```

In [0]: # LEARNING CURVE SAMPLES UNITS

```

def plot_learning_curve_samples(estimator, title, X, y, ylim=None, cv=None,
                                n_jobs=None, train_sizes=np.linspace(.1, 1.0, 5)):
    plt.figure()
    plt.title(title)
    if ylim is not None:
        plt.ylim(*ylim)
    plt.xlabel("Training examples")

```

```

plt.ylabel("Score")
train_sizes, train_scores, test_scores = learning_curve(
    estimator, X, y, cv=cv, n_jobs=n_jobs, train_sizes=train_sizes)
train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)
test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)
plt.grid()

plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                 train_scores_mean + train_scores_std, alpha=0.1,
                 color="r")
plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                 test_scores_mean + test_scores_std, alpha=0.1, color="g")
plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
         label="Training score")
plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
         label="Validation score")

plt.legend(loc="best")
return plt

```

In [0]: # VALIDATION CURVE only plotting

```

def plot_validation_curve(train_scores, test_scores, param_range):
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)
    plt.title("Validation Curve")
    plt.xlabel(r"$\gamma$")
    plt.ylabel("Score")
    #plt.ylim(0.0, 1.1)
    lw = 2
    plt.semilogx(param_range, train_scores_mean, label="Training score",
                 color="darkorange", lw=lw)
    plt.fill_between(param_range, train_scores_mean - train_scores_std,
                     train_scores_mean + train_scores_std, alpha=0.2,
                     color="darkorange", lw=lw)
    plt.semilogx(param_range, test_scores_mean, label="Cross-validation score",
                 color="navy", lw=lw)
    plt.fill_between(param_range, test_scores_mean - test_scores_std,
                     test_scores_mean + test_scores_std, alpha=0.2,
                     color="navy", lw=lw)
    plt.legend(loc="best")
    plt.show()

```

In [0]: # CROSS VALIDATION SCORE

```
def display_scores(scores):
    print('RMSE_Scores: ',scores)
    print("RMSE_Scores_Mean:", scores.mean())
    print("RMSE_Scores_Standard deviation:", scores.std())
```

In [0]: # CROSS VALIDATION SCORE - CALCULATION AND OUTPUT

```
def cross_val_score_full(estimator,X,y,scoring,cv):
    scores=cross_val_score(estimator,X,y,scoring,cv)
    scores_rmse=np.sqrt(-scores)
    print('Scores: ',scores_rmse)
    print("Mean:", scores_rmse.mean())
    print("Standard deviation:", scores_rmse.std())
```

## 9 ML REGRESSION ALGORITHMS MODELS

1. SGD REGRESSOR
2. LINEAR REGRESSION
3. LINEAR - SVR
4. SVR
5. DECISION TREE
6. KNEIGHBORS 7.RANDOMFOREST
7. BAGGING
8. VOTING
9. REGULARIZATION
10. BOOSTING GRADIENT REGRESSOR
11. NEURAL NETWORK

### Models

- Original Data with Scaling and conversion of categorical data
- Data with added Polynomial Features

### Performance Metrics Methodologies:

- Regression Metrics - R2/RMSE/Absolute Error
- Cross Validation Score - Mean of R2
- Learning Curve - R2 and evaluation based on # of various samples bins i.e. 100,1000,10000
- Validation Curve - R2 and evaluation based on hyper parameters

The Mean Absolute Error (or MAE) is the average of the absolute differences between predictions and actual values. It gives an idea of how wrong the predictions were. The Mean Squared Error (or MSE) is much like the mean absolute error in that it provides a gross idea of the magnitude of error. Taking the square root of the mean squared error converts the units back to the original units of the output variable and can be meaningful for description and presentation. This is called the Root Mean Squared Error (or RMSE). The  $R^2$  (or R Squared) metric provides an indication of the goodness of fit of a set of predictions to the actual values. In statistical literature, this measure is called the coefficient of determination.

```
In [0]: # Variable Set up
```

```
cv=10
```

### #1.SGD-REGRESSOR EVALUATION

Following steps followed to develop models:

First train the data with algorithm and checked overfitting or not through cross validation and learning curve method Second if the model is not overfitting then applied feature engineering through polynomial feature and checked overfitting or not through cross validation and learning curve method

Lastly, calculated best parameters related to the model

```
In [28]: # Basic Set up related to SGDRegressor
```

```
from sklearn.linear_model import SGDRegressor
lin_reg_SGD=SGDRegressor(random_state=42)
lin_reg_SGD.get_params()
```

```
Out[28]: {'alpha': 0.0001,
          'average': False,
          'early_stopping': False,
          'epsilon': 0.1,
          'eta0': 0.01,
          'fit_intercept': True,
          'l1_ratio': 0.15,
          'learning_rate': 'invscaling',
          'loss': 'squared_loss',
          'max_iter': 1000,
          'n_iter_no_change': 5,
          'penalty': 'l2',
          'power_t': 0.25,
          'random_state': 42,
          'shuffle': True,
          'tol': 0.001,
          'validation_fraction': 0.1,
          'verbose': 0,
          'warm_start': False}
```

```
In [0]: # Train the data with no added features with performance metrics
```

```
a=lin_reg_SGD.fit(X_train,y_train)
y_pred_SGD=a.predict(X_train)

lin_perform_metrics(y_train,y_pred_SGD)
```

```
R2_Square: 0.6327407931679956
Explained_Variance_Score: 0.6331968948040358
Mean_absolute_error: 79.65184408900929
Median_absolute_error: 57.46733435379285
```

Max\_error: 474.205753868473  
MSE: 12033.309543955203  
RMSE: 109.69644271331319

Observations: R2 score is low which means this model prediction is not good which is also evident from RMSE high score when you consider y\_train mean is 189.

```
In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2
        cross_val_score(lin_reg_SGD,X_train,y_train,cv=cv).mean()
```

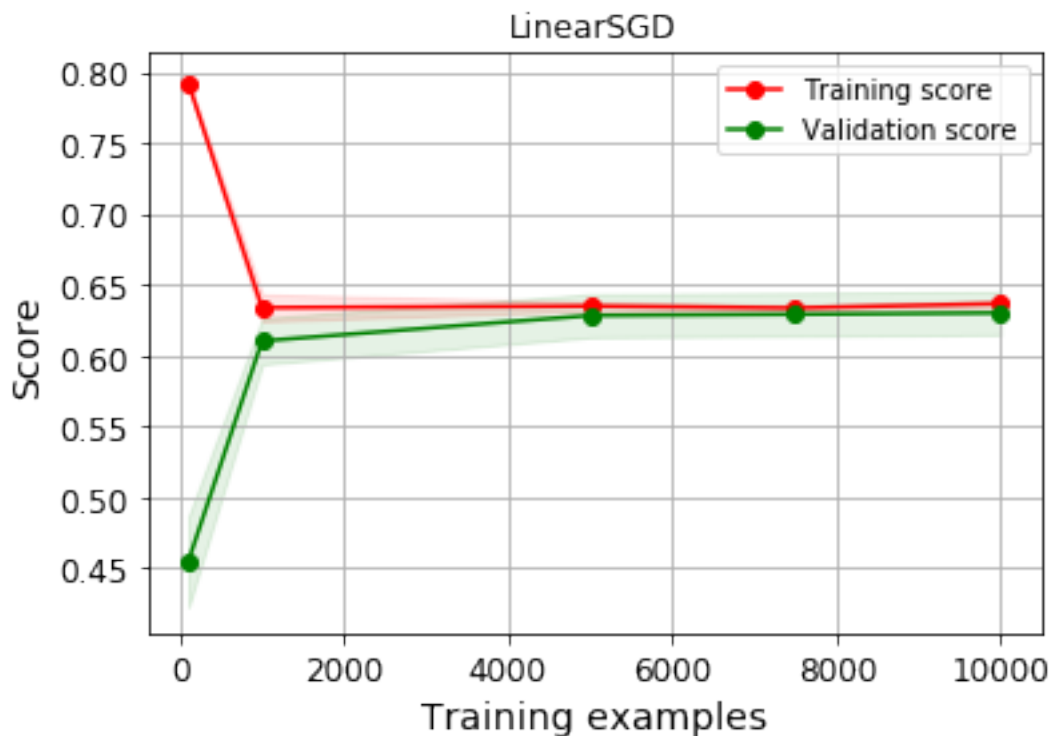
Out[0]: 0.6300484658017598

Observations: CV R2 score is very close to the original result which means that validation model is performing well and it is not overfitting the dataset.

```
In [32]: # LEARNING CURVE TO EVALUATE PERFORMANCE BASED ON BLOCKS OF SAMPLES
```

```
size=[100,1000,5000,7500,10000]
plot_learning_curve_samples(lin_reg_SGD, 'LinearSGD', X_train, y_train, ylim=None, cv=
                             n_jobs=-1, train_sizes=size)
```

Out[32]: <module 'matplotlib.pyplot' from '/usr/local/lib/python3.6/dist-packages/matplotlib/p



Observations: Learning curve result is showing that this model best R2 score are flat after 5000 samples .

```
In [0]: # SGD model applying on Train Data with feature engineering
```

```
lin_reg_SGD.fit(X_poly,y_train)
y_pred_SGD_poly=lin_reg_SGD.predict(X_poly)
lin_perform_metrics(y_train,y_pred_SGD_poly)
```

```
R2_Square: 0.8549733499424739
Explained_Variance_Score: 0.8549906866471491
Mean_absolute_error: 50.238714902851164
Median_absolute_error: 36.237785810575645
Max_error: 406.9391573081589
MSE: 4751.822526979878
RMSE: 68.93346449279825
```

The model performance has improved over 32% just adding one degree of feature.

```
In [0]: # Evaluate Cross Val Score on Dataset with added features
```

```
cross_val_score(lin_reg_SGD,X_poly,y_train,cv=cv).mean()
```

```
Out[0]: 0.8370798032303745
```


Observation: CV score is lower than the model without CV.

```
In [33]: plot_learning_curve_samples(lin_reg_SGD, 'LinearSGD_additional_features', X_poly, y_train,
```

```
                                     n_jobs=-1, train_sizes=size)
```

```
save_results_to = ''
```

```
plt.savefig(save_results_to + 'LinearSGD', dpi = 300)
```





Observation: Validation score is close to the training model which it is not overfitted yet.

```
In [0]: # DETERMINE BEST PARAMETERS RESULTED ASSOCIATED WITH X_POLY DATA
alpha=[0.0001,0.001,0.01,0.1,1,10,100]
penalty=['l2','l1','elasticnet']
param ={'alpha':alpha,'penalty': penalty}
best_parameters_poly(lin_reg_SGD,param,cv=cv)
```

Fitting 10 folds for each of 21 candidates, totalling 210 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 46 tasks      | elapsed: 23.5s
[Parallel(n_jobs=-1)]: Done 196 tasks    | elapsed: 1.4min
[Parallel(n_jobs=-1)]: Done 210 out of 210 | elapsed: 1.6min finished
```

Best\_results: 0.837694679270226

Best\_parameters: {'alpha': 0.01, 'penalty': 'l1'}

Best\_estimator: SGDRegressor(alpha=0.01, average=False, early\_stopping=False, epsilon=0.1, eta0=0.01, fit\_intercept=True, l1\_ratio=0.15, learning\_rate='invscaling', loss='squared\_loss', max\_iter=1000, n\_iter\_no\_change=5, penalty='l1', power\_t=0.25, random\_state=42, shuffle=True, tol=0.001, validation\_fraction=0.1, verbose=0, warm\_start=False)

**CONCLUSION:** SGD REGRESSOR MODEL WITH 2 DEGREE OF POLYNOMIAL IS PERFORMING VERY WELL WITH R2 SCORE IS ABOVE 0.80.

#2.LINEAR REGRESSOR EVALUATION

```
In [34]: from sklearn.linear_model import LinearRegression
lin_reg=LinearRegression()
lin_reg.get_params()
```

```
Out[34]: {'copy_X': True, 'fit_intercept': True, 'n_jobs': None, 'normalize': False}
```

```
In [0]: # Train the data with no added features with performance metrics with no CV
```

```
lin_reg.fit(X_train,y_train)
y_pred_lin_reg=lin_reg.predict(X_train)

lin_perform_metrics(y_train,y_pred_lin_reg)
```

R2\_Square: 0.6333171267805959

Explained\_Variance\_Score: 0.6333171267805959

Mean\_absolute\_error: 79.25944079175646

Median\_absolute\_error: 56.654444408712266  
Max\_error: 476.38000751926825  
MSE: 12014.425876420142  
RMSE: 109.61033654003687

Observations: R2 score is low which means this model prediction is not good which is also evident from RMSE high score when you consider y\_train mean is 189.

In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2

```
cross_val_score(lin_reg,X_train,y_train,cv=cv).mean()
```

Out[0]: 0.6301644002256567

Observations: CV R2 score is very close to the original result which means that validation model is performing well and it is not overfitting the dataset.

In [0]: # LEARNING CURVE TO EVALUATE PERFORMANCE BASED ON BLOCKS OF SAMPLES

```
plot_learning_curve_samples(lin_reg, 'LinearRegression', X_train, y_train, ylim=None,
                             n_jobs=-1, train_sizes=size)
```

Out[0]: <module 'matplotlib.pyplot' from '/usr/local/lib/python3.6/dist-packages/matplotlib/py...



Comments: Learning curve result is showing that this model best R2 score are flat after 5000 samples .

```
In [0]: # Linear model applying on Train Data with feature engineering
```

```
lin_reg.fit(X_poly,y_train)
y_pred_linreg_poly=lin_reg.predict(X_poly)
lin_perform_metrics(y_train,y_pred_linreg_poly)
```

```
R2_Square: 0.8587717894563467
Explained_Variance_Score: 0.8587717894563467
Mean_absolute_error: 49.499634026324465
Median_absolute_error: 35.79412952745116
Max_error: 410.5645089519726
MSE: 4627.366018867527
RMSE: 68.02474563618395
```

The model performance has improved by 37% just adding one degree of feature.

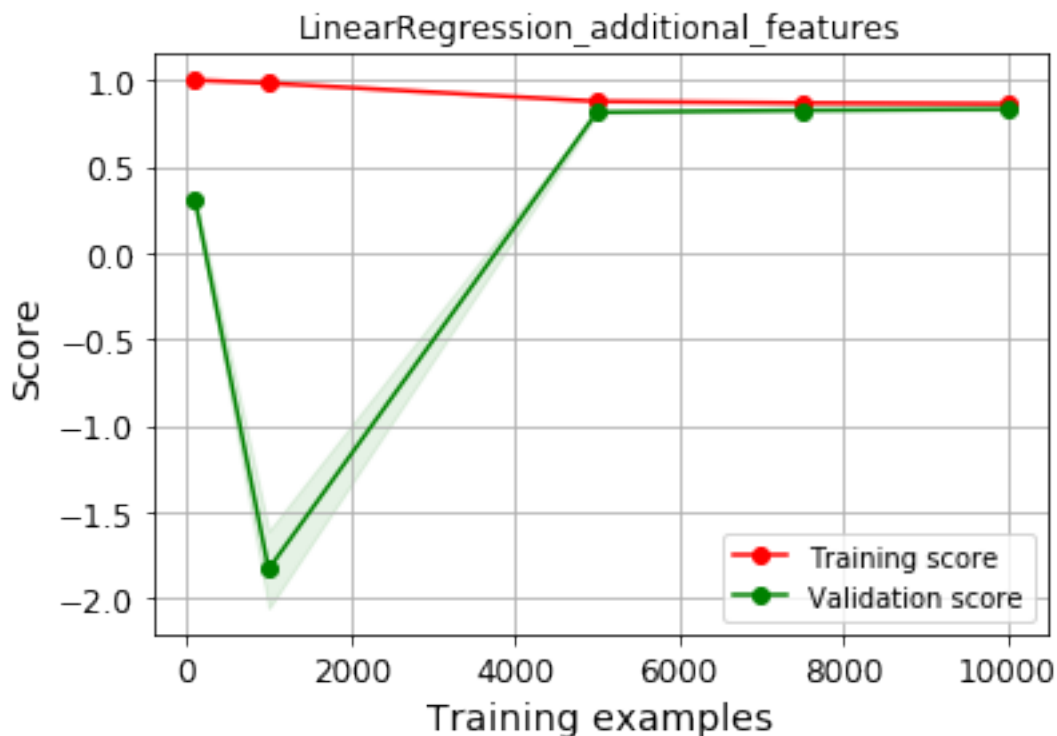
```
In [0]: # Evaluate Cross Val Score on Dataset with added features
```

```
cross_val_score(lin_reg,X_poly,y_train,cv=cv).mean()
```

```
Out[0]: 0.8359940726145121
```

Observation: CV score is lower than the model without CV.

```
In [35]: plot_learning_curve_samples(lin_reg, 'LinearRegression_additional_features', X_poly, y_train,
                                     n_jobs=-1, train_sizes=size)
save_results_to = ''
plt.savefig(save_results_to + 'Linear_Regression', dpi = 300)
```



Observation: Validation score is close to the training model which it is not overfitted yet.

**CONCLUSION:** LINEAR REGRESSOR MODEL WITH 2 DEGREE OF POLYNOMIAL IS PERFORMING VERY WELL WITH R2 SCORE IS ABOVE 0.80.

## 10 3.LINEAR-SVR EVALUATION

```
In [0]: from sklearn.svm import LinearSVR
lin_SVR=LinearSVR(random_state=42)
lin_SVR.get_params
```

```
Out[0]: <bound method BaseEstimator.get_params of LinearSVR(C=1.0, dual=True, epsilon=0.0, fit_intercept_scaling=1.0, loss='epsilon_insensitive', max_iter=1000, random_state=42, tol=0.0001, verbose=0)>
```

```
In [0]: # Train the data with no added features with performance metrics with no CV
```

```
lin_SVR.fit(X_train,y_train)
y_pred_lin_SVR=lin_SVR.predict(X_train)

lin_perform_metrics(y_train,y_pred_lin_SVR)
```


```
R2_Square: 0.5754417921712207
Explained_Variance_Score: 0.5901634118536487
Mean_absolute_error: 79.24870885835276
Median_absolute_error: 50.77158249940891
Max_error: 595.2125212770032
MSE: 13908.940601622851
RMSE: 117.93617172701025
```

```
In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2
cross_val_score(lin_SVR,X_train,y_train,cv=cv).mean()
```

```
Out[0]: 0.5676880268073471
```

```
In [0]: # LEARNING CURVE TO EVALUATE PERFORMANCE BASED ON BLOCKS OF SAMPLES
plot_learning_curve_samples(lin_SVR, 'LinearSVR', X_train, y_train, ylim=None, cv=cv,
                             n_jobs=-1, train_sizes=size)
```

```
Out[0]: <module 'matplotlib.pyplot' from '/usr/local/lib/python3.6/dist-packages/matplotlib/py
```



In [0]: # model applying on Train Data with feature engineering

```
lin_SVR.fit(X_poly,y_train)
y_pred_lin_SVR_poly=lin_SVR.predict(X_poly)
lin_perform_metrics(y_train,y_pred_lin_SVR_poly)
```

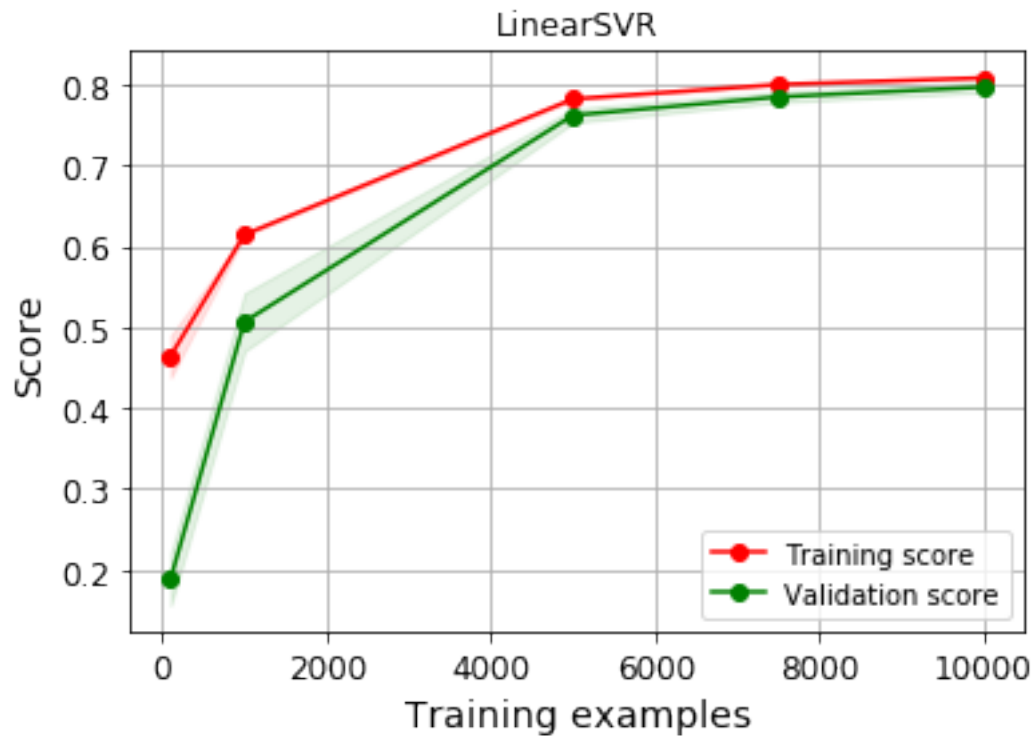
R2\_Square: 0.8162697337062798  
 Explained\_Variance\_Score: 0.8196041501312942  
 Mean\_absolute\_error: 49.63861954361137  
 Median\_absolute\_error: 28.63086978139694  
 Max\_error: 440.39504984059295  
 MSE: 6019.182560781661  
 RMSE: 77.58339101110276

In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2  
 cross\_val\_score(lin\_SVR,X\_poly,y\_train,cv=cv).mean()

Out[0]: 0.8038758777181941

In [0]: # LEARNING CURVE TO EVALUATE PERFORMANCE BASED ON BLOCKS OF SAMPLES  
 plot\_learning\_curve\_samples(lin\_SVR, 'LinearSVR', X\_poly, y\_train, ylim=None, cv=cv,  
 n\_jobs=-1, train\_sizes=size)

```
Out[0]: <module 'matplotlib.pyplot' from '/usr/local/lib/python3.6/dist-packages/matplotlib/py
```



```
In [0]: # DETERMINE BEST PARAMETERS RESULTED ASSOCIATED WITH X_POLY DATA
param ={'C':[0.3,0.6,1,10,100],'epsilon': [0.01,0.03,0.06,1,10,100]}
best_parameters_poly(lin_SVR,param,cv=cv)
```

Fitting 10 folds for each of 30 candidates, totalling 300 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 88 tasks      | elapsed:    5.9s
[Parallel(n_jobs=-1)]: Done 289 tasks    | elapsed:   3.8min
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed:   4.1min finished
```

```
Best_results: 0.824137422871571
Best_parameters: {'C': 10, 'epsilon': 10}
Best_estimator: LinearSVR(C=10, dual=True, epsilon=10, fit_intercept=True,
    intercept_scaling=1.0, loss='epsilon_insensitive', max_iter=1000,
    random_state=42, tol=0.0001, verbose=0)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:929: ConvergenceWarning: Liblinear :
    "the number of iterations.", ConvergenceWarning)
```

## CONCLUSION:

Summary: Dataset with feature engineering resulted in  $R^2$  of 0.8067 for Linear SVR algorithm.

## #4.SVR EVALUATION

```
In [0]: from sklearn.svm import SVR
        SVR=SVR()
        SVR.get_params()
```

```
Out[0]: {'C': 1.0,
         'cache_size': 200,
         'coef0': 0.0,
         'degree': 3,
         'epsilon': 0.1,
         'gamma': 'auto_deprecated',
         'kernel': 'rbf',
         'max_iter': -1,
         'shrinking': True,
         'tol': 0.001,
         'verbose': False}
```

```
In [0]: # SVR model applying on Train Data WITHOUT feature engineering
```

```
SVR.fit(X_train,y_train)
y_pred_SVR=SVR.predict(X_train)
lin_perform_metrics(y_train,y_pred_SVR)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will change from 'auto_deprecated' to 'scale' in version 0.22. To avoid this warning, you can explicitly set gamma to 'scale' or 'auto'.
"avoid this warning.", FutureWarning)
```

```
R2_Square: 0.2611777119405837
Explained_Variance_Score: 0.3114429377267419
Mean_absolute_error: 106.4164886633958
Median_absolute_error: 71.39670635470401
Max_error: 737.308768569687
MSE: 24204.538106392774
RMSE: 155.57807720367538
```

```
In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2
        cross_val_score(SVR,X_train,y_train,cv=cv).mean()
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will change from 'auto_deprecated' to 'scale' in version 0.22. To avoid this warning, you can explicitly set gamma to 'scale' or 'auto'.
"avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will change from 'auto_deprecated' to 'scale' in version 0.22. To avoid this warning, you can explicitly set gamma to 'scale' or 'auto'.
"avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will change from 'auto_deprecated' to 'scale' in version 0.22. To avoid this warning, you can explicitly set gamma to 'scale' or 'auto'.
"avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will change from 'auto_deprecated' to 'scale' in version 0.22. To avoid this warning, you can explicitly set gamma to 'scale' or 'auto'.
"avoid this warning.", FutureWarning)
```

```

"avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will be 'auto' in the future.
"avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will be 'auto' in the future.
"avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will be 'auto' in the future.
"avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will be 'auto' in the future.
"avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will be 'auto' in the future.
"avoid this warning.", FutureWarning)

```

```
Out[0]: 0.2525148695261212
```

```

In [0]: # learning curve evaluation
        plot_learning_curve_samples(SVR, 'SVR', X_train, y_train, ylim=None, cv=cv,
                                   n_jobs=-1, train_sizes=size)

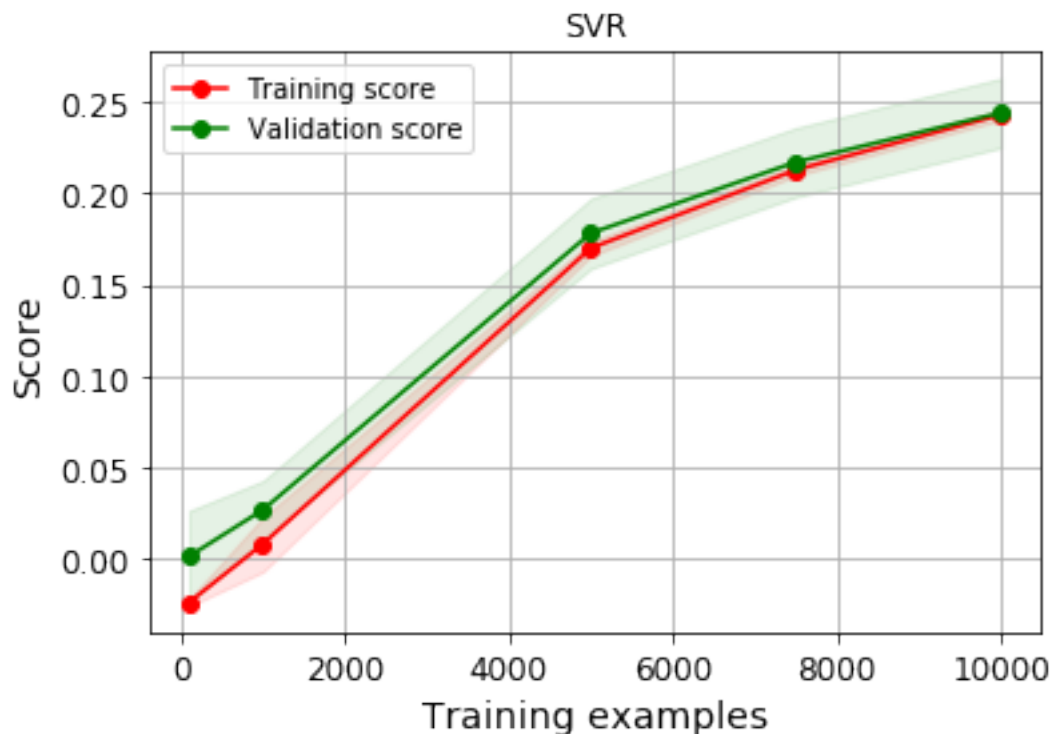
```

```

/usr/local/lib/python3.6/dist-packages/joblib/externals/loky/process_executor.py:706: UserWarning:
"timeout or by a memory leak.", UserWarning

```

```
Out[0]: <module 'matplotlib.pyplot' from '/usr/local/lib/python3.6/dist-packages/matplotlib/pyplot.py'>
```





```
In [0]: # SVR model applying on Train Data WITH feature engineering
```

```
SVR.fit(X_poly,y_train)
y_pred_SVR_poly=SVR.predict(X_poly)
lin_perform_metrics(y_train,y_pred_SVR_poly)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22.
  "avoid this warning.", FutureWarning)
```

```
R2_Square: 0.10168956079933955
Explained_Variance_Score: 0.16798238259153275
Mean_absolute_error: 120.7601091775984
Median_absolute_error: 88.74450057622401
Max_error: 790.815318289249
MSE: 29253.702856753232
RMSE: 171.03713882298555
```

```
In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2
```

```
cross_val_score(SVR,X_poly,y_train,cv=cv).mean()
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22.
  "avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22.
  "avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22.
  "avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22.
  "avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22.
  "avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22.
  "avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22.
  "avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22.
  "avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22.
  "avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22.
  "avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22.
  "avoid this warning.", FutureWarning)
```

```
Out[0]: 0.0910090887185098
```

```
In [0]: param ={'kernel':['linear', 'poly', 'rbf']}
        best_parameters_poly(SVR,param,cv=cv)
```

Fitting 10 folds for each of 3 candidates, totalling 30 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 30 out of 30 | elapsed: 20.6min finished
```

```
Best_results: 0.8035877593540758
Best_parameters: {'kernel': 'linear'}
Best_estimator: SVR(C=1.0, cache_size=200, coef0=0.0, degree=3, epsilon=0.1,
    gamma='auto_deprecated', kernel='linear', max_iter=-1, shrinking=True,
    tol=0.001, verbose=False)
```

Summary: Dataset with feature engineering resulted in  $R^2$  of 0.8036 for SV Regressor algorithm.

#### #5.DECISIONTREEREgressor

```
In [0]: from sklearn.tree import DecisionTreeRegressor
        decision_tree_reg=DecisionTreeRegressor()
        decision_tree_reg.get_params
```

```
Out[0]: <bound method BaseEstimator.get_params of DecisionTreeRegressor(criterion='mse', max_depth=
        max_leaf_nodes=None, min_impurity_decrease=0.0,
        min_impurity_split=None, min_samples_leaf=1,
        min_samples_split=2, min_weight_fraction_leaf=0.0,
        presort=False, random_state=None, splitter='best')>
```

```
In [0]: # Model applying on Train Data WITHOUT feature engineering
```

```
        decision_tree_reg.fit(X_train,y_train)
        y_pred_DT=decision_tree_reg.predict(X_train)
        lin_perform_metrics(y_train,y_pred_DT)
```

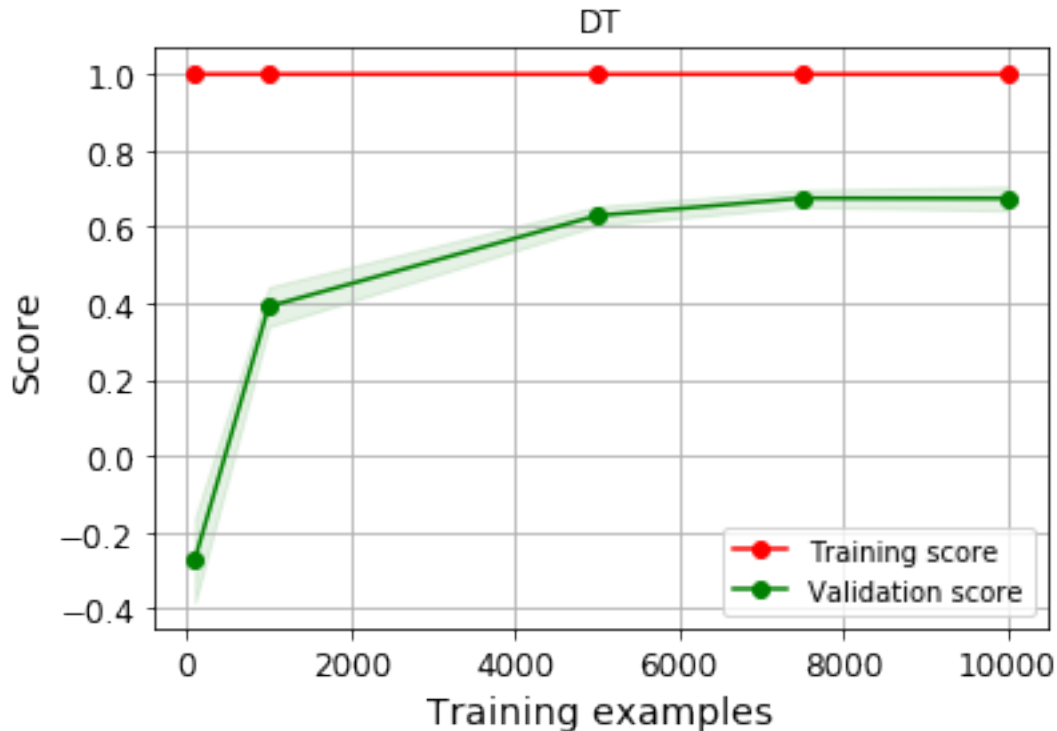
```
R2_Square: 0.9999628827723317
Explained_Variance_Score: 0.9999628827723317
Mean_absolute_error: 0.030640868877220742
Median_absolute_error: 0.0
Max_error: 67.0
MSE: 1.2159965475077321
RMSE: 1.1027223347278916
```

```
In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2
        cross_val_score(decision_tree_reg,X_train,y_train,cv=cv).mean()
```

```
Out[0]: 0.6928397674269655
```

```
In [0]: # learning curve evaluation
        plot_learning_curve_samples(decision_tree_reg, 'DT', X_train, y_train, ylim=None, cv=cv,
                                   n_jobs=-1, train_sizes=size)
```

```
Out[0]: <module 'matplotlib.pyplot' from '/usr/local/lib/python3.6/dist-packages/matplotlib/pyplot.py'>
```



Comments: Not performing well in cv due to overfitting.

```
In [0]: # model applying on Train Data WITH feature engineering
```

```
decision_tree_reg.fit(X_poly,y_train)
y_pred_DT_poly=decision_tree_reg.predict(X_poly)
lin_perform_metrics(y_train,y_pred_DT_poly)
```

```
R2_Square: 0.9999628827723317
Explained_Variance_Score: 0.9999628827723317
Mean_absolute_error: 0.030640868877220742
Median_absolute_error: 0.0
Max_error: 67.0
MSE: 1.2159965475077321
RMSE: 1.1027223347278916
```

```
In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2
        cross_val_score(decision_tree_reg,X_poly,y_train,cv=cv).mean()
```

```
Out[0]: 0.6850356155932438
```

```
In [0]: # DETERMINE BEST PARAMETERS RESULTED ASSOCIATED WITH X_TRAIN DATA
```

```
param={'max_features': [0.69,0.7,0.71], 'max_depth':[10,12,15], 'min_samples_leaf':[0.001,0.01,0.1]}
```

```
best_parameters(decision_tree_reg,param,cv=cv)
```

Fitting 10 folds for each of 27 candidates, totalling 270 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 88 tasks      | elapsed:      4.6s
```

```
Best_results: 0.7254413126025471
```

```
Best_parameters: {'max_depth': 15, 'max_features': 0.71, 'min_samples_leaf': 0.001}
```

```
Best_estimator: DecisionTreeRegressor(criterion='mse', max_depth=15, max_features=0.71,
                                     max_leaf_nodes=None, min_impurity_decrease=0.0,
                                     min_impurity_split=None, min_samples_leaf=0.001,
                                     min_samples_split=2, min_weight_fraction_leaf=0.0,
                                     presort=False, random_state=None, splitter='best')
```

```
[Parallel(n_jobs=-1)]: Done 270 out of 270 | elapsed:    19.4s finished
```

Summary: Dataset without feature engineering resulted in  $R^2$  of 0.7254 for Decision Tree Regressor algorithm.

#### #6.KNEIGHBORS-REGRESSOR

```
In [0]: from sklearn.neighbors import KNeighborsRegressor
```

```
knn_reg=KNeighborsRegressor()
```

```
knn_reg.get_params
```

```
Out[0]: <bound method BaseEstimator.get_params of KNeighborsRegressor(algorithm='auto', leaf_size=30,
                                metric_params=None, n_jobs=None, n_neighbors=5, p=2,
                                weights='uniform')>
```

```
In [0]: # Model applying on Train Data WITHOUT feature engineering
```

```
knn_reg.fit(X_train,y_train)
```

```
y_pred_knn=knn_reg.predict(X_train)
```

```
lin_perform_metrics(y_train,y_pred_knn)
```

```
R2_Square: 0.7541366768751541
```

```
Explained_Variance_Score: 0.7541588787900734
```

```
Mean_absolute_error: 60.90685463569014
```

```
Median_absolute_error: 40.0
```

```
Max_error: 616.8
MSE: 8054.722048478745
RMSE: 89.74810331410211
```

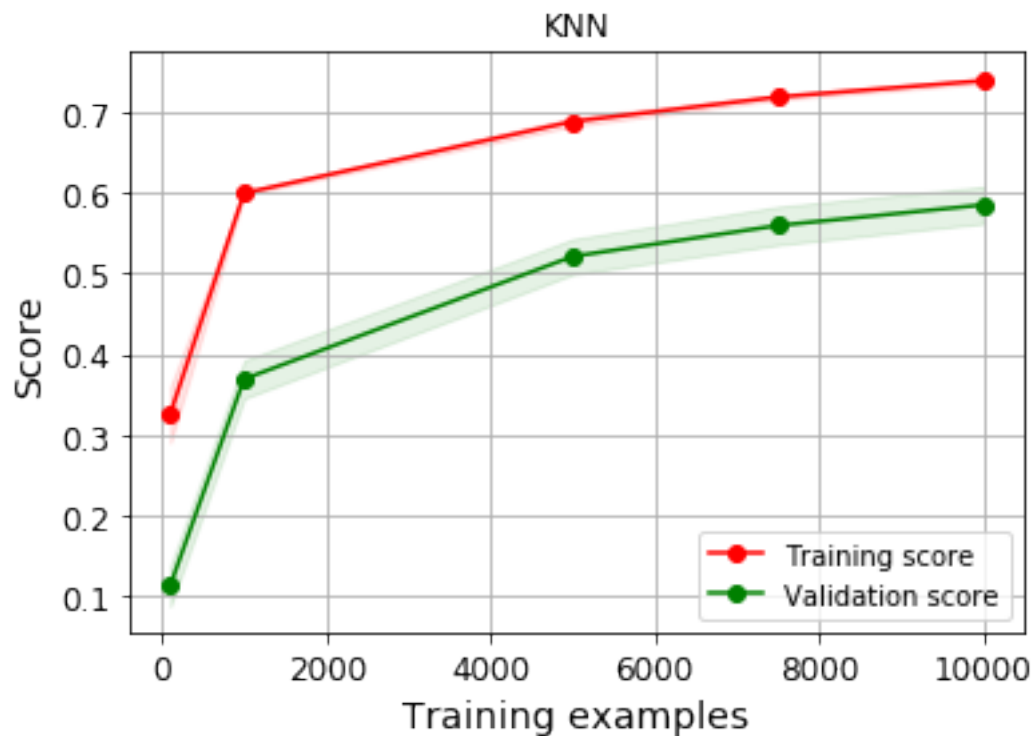
```
In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2
        cross_val_score(knn_reg,X_train,y_train,cv=cv).mean()
```

```
Out[0]: 0.6080392303948567
```

```
In [0]: # learning curve evaluation
        plot_learning_curve_samples(knn_reg, 'KNN', X_train, y_train, ylim=None, cv=cv,
                                     n_jobs=-1, train_sizes=size)
```

```
/usr/local/lib/python3.6/dist-packages/joblib/externals/loky/process_executor.py:706: UserWarning:
  "timeout or by a memory leak.", UserWarning
```

```
Out[0]: <module 'matplotlib.pyplot' from '/usr/local/lib/python3.6/dist-packages/matplotlib/py...
```



```
In [0]: knn_reg.fit(X_poly,y_train)
        y_pred_knn_poly=knn_reg.predict(X_poly)
        lin_perform_metrics(y_train,y_pred_knn_poly)
```

```
R2_Square: 0.7570040774874776
Explained_Variance_Score: 0.7570081482212435
Mean_absolute_error: 61.08947709127526
Median_absolute_error: 40.19999999999999
Max_error: 576.6
MSE: 7996.609126087895
RMSE: 89.42376152951684
```

```
In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2
        cross_val_score(knn_reg,X_poly,y_train,cv=cv).mean()
```

```
Out[0]: 0.6047874076922802
```

```
In [0]: # DETERMINE BEST PARAMETERS RESULTED ASSOCIATED WITH X_TRAIN DATA
        param=[{'n_neighbors': [5,10,15], 'leaf_size': [30,40,50],}]
        best_parameters_poly(knn_reg,param,cv=cv)
```

Fitting 10 folds for each of 9 candidates, totalling 90 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 46 tasks      | elapsed: 1.3min
```

```
Best_results: 0.5946768752989717
Best_parameters: {'leaf_size': 30, 'n_neighbors': 5}
Best_estimator: KNeighborsRegressor(algorithm='auto', leaf_size=30, metric='minkowski',
                                     metric_params=None, n_jobs=None, n_neighbors=5, p=2,
                                     weights='uniform')
```

```
[Parallel(n_jobs=-1)]: Done 90 out of 90 | elapsed: 2.5min finished
```

Summary: Dataset with feature engineering resulted in  $R^2$  of 0.5946 for KNeighbors Regressor algorithm.

#### #7.RANDOMFOREST - REGRESSOR

```
In [36]: from sklearn.ensemble import RandomForestRegressor
        random_reg=RandomForestRegressor()
        random_reg.get_params()
```

```
Out[36]: {'bootstrap': True,
          'criterion': 'mse',
          'max_depth': None,
          'max_features': 'auto',
          'max_leaf_nodes': None,
          'min_impurity_decrease': 0.0,
```

```
'min_impurity_split': None,  
'min_samples_leaf': 1,  
'min_samples_split': 2,  
'min_weight_fraction_leaf': 0.0,  
'n_estimators': 'warn',  
'n_jobs': None,  
'oob_score': False,  
'random_state': None,  
'verbose': 0,  
'warm_start': False}
```

```
In [0]: random_reg.fit(X_train,y_train)  
        y_pred_random=random_reg.predict(X_train)  
        lin_perform_metrics(y_train,y_pred_random)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The default  
"10 in version 0.20 to 100 in 0.22.", FutureWarning)
```


```
R2_Square: 0.9671568885815681  
Explained_Variance_Score: 0.967193978160988  
Mean_absolute_error: 19.651861588625955  
Median_absolute_error: 10.100000000000023  
Max_error: 307.7  
MSE: 1076.1100572364478  
RMSE: 32.804116467852744
```

```
In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2  
        cross_val_score(random_reg,X_train,y_train,cv=cv).mean()
```

```
Out[0]: 0.8202394719092576
```

```
In [0]: # learning curve evaluation  
        plot_learning_curve_samples(random_reg, 'Random_Forest', X_train, y_train, ylim=None, c  
                                     n_jobs=-1, train_sizes=size)
```

```
Out[0]: <module 'matplotlib.pyplot' from '/usr/local/lib/python3.6/dist-packages/matplotlib/py
```



```
In [0]: # from Hithendar
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import GridSearchCV

regressor = RandomForestRegressor()
parameters = [{'n_estimators' : [150,200,250,300], 'max_features' : ['auto','sqrt','log2']}
grid_search = GridSearchCV(estimator = regressor, param_grid = parameters)
grid_search = grid_search.fit(X_train, y_train)
best_parameters = grid_search.best_params_
best_accuracy = grid_search.best_score_

/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_split.py:1978: FutureWarning:
  warnings.warn(CV_WARNING, FutureWarning)

In [0]: print(best_parameters)
print(best_accuracy)

{'max_features': 'auto', 'n_estimators': 300}
0.8312165428200716

In [0]: param = [{'n_estimators' : [150,200,250,300], 'max_features' : ['auto','sqrt','log2']}
best_parameters(random_reg,param,cv=2)
# this is very time consuming process....
```



Fitting 2 folds for each of 12 candidates, totalling 24 fits


```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.  
/usr/local/lib/python3.6/dist-packages/joblib/externals/loky/process_executor.py:706: UserWarning:  
    "timeout or by a memory leak.", UserWarning  
[Parallel(n_jobs=-1)]: Done 24 out of 24 | elapsed: 8.2min finished
```

```
Best_results: 0.8262559668916158  
Best_parameters: {'max_features': 'auto', 'n_estimators': 300}  
Best_estimator: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,  
    max_features='auto', max_leaf_nodes=None,  
    min_impurity_decrease=0.0, min_impurity_split=None,  
    min_samples_leaf=1, min_samples_split=2,  
    min_weight_fraction_leaf=0.0, n_estimators=300,  
    n_jobs=None, oob_score=False, random_state=None,  
    verbose=0, warm_start=False)
```

```
In [0]: random_reg.fit(X_poly,y_train)  
        y_pred_random_poly=random_reg.predict(X_poly)  
        lin_perform_metrics(y_train,y_pred_random_poly)
```

```
R2_Square: 0.9678316387287264  
Explained_Variance_Score: 0.9678319020423837  
Mean_absolute_error: 19.68017250473519  
Median_absolute_error: 10.299999999999997  
Max_error: 413.4  
MSE: 1054.0017554306905  
RMSE: 32.4653931969211
```

```
In [37]: # learning curve evaluation  
        plot_learning_curve_samples(random_reg, 'Random_Forest', X_poly, y_train, ylim=None, c  
            n_jobs=-1, train_sizes=size)  
        save_results_to = ''  
        plt.savefig(save_results_to + 'Random_Forest_additional_features', dpi = 300)
```



```
In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2
        cross_val_score(random_reg,X_poly,y_train,cv=cv).mean()
```

```
Out[0]: 0.8203102297919515
```

Summary: Dataset without feature engineering resulted in  $R^2$  of 0.8337 for Random Forest algorithm.

#8.BAGGING -REGRESSOR

```
In [0]: from sklearn.ensemble import BaggingRegressor
        bag_reg=BaggingRegressor(random_state=42)
        bag_reg.get_params()
```

```
Out[0]: {'base_estimator': None,
         'bootstrap': True,
         'bootstrap_features': False,
         'max_features': 1.0,
         'max_samples': 1.0,
         'n_estimators': 10,
         'n_jobs': None,
         'oob_score': False,
         'random_state': 42,
         'verbose': 0,
         'warm_start': False}
```

```
In [0]: bag_reg.fit(X_train,y_train)
        y_pred_bag=bag_reg.predict(X_train)
        lin_perform_metrics(y_train,y_pred_bag)
```


```
R2_Square: 0.9693556518551601
Explained_Variance_Score: 0.9693754572621
Mean_absolute_error: 19.3224519887794
Median_absolute_error: 9.900000000000006
Max_error: 264.5
MSE: 1008.4567325421373
RMSE: 31.7562077796159
```

```
In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2
        cross_val_score(bag_reg,X_train,y_train,cv=cv).mean()
```

```
Out[0]: 0.8215451437411845
```

```
In [0]: # learning curve evaluation
        plot_learning_curve_samples(bag_reg, 'Bagging', X_train, y_train, ylim=None, cv=cv,
                                     n_jobs=-1, train_sizes=size)
```

```
Out[0]: <module 'matplotlib.pyplot' from '/usr/local/lib/python3.6/dist-packages/matplotlib/py...
```



```
In [0]: bag_reg.fit(X_poly,y_train)
        y_pred_bag_poly=bag_reg.predict(X_poly)
        lin_perform_metrics(y_train,y_pred_bag_poly)
```

```
R2_Square: 0.9681550747459364
Explained_Variance_Score: 0.9681570620094164
Mean_absolute_error: 19.612846627826126
Median_absolute_error: 10.099999999999994
Max_error: 313.79999999999995
MSE: 1047.965814706658
RMSE: 32.37230011455253
```

```
In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2
        cross_val_score(bag_reg,X_poly,y_train,cv=cv).mean()
```

```
Out[0]: 0.821024962942753
```

```
In [0]: # DETERMINE BEST PARAMETERS RESULTED ASSOCIATED WITH X_TRAIN DATA
        param={'max_features': [1], 'n_estimators': [10]}
        best_parameters_poly(bag_reg,param,cv=cv)
```

Fitting 10 folds for each of 1 candidates, totalling 10 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.

```
Best_results: 0.009712246171741722
Best_parameters: {'max_features': 1, 'n_estimators': 10}
Best_estimator: BaggingRegressor(base_estimator=None, bootstrap=True, bootstrap_features=False,
                                max_features=1, max_samples=1.0, n_estimators=10, n_jobs=None,
                                oob_score=False, random_state=42, verbose=0, warm_start=False)
```

[Parallel(n\_jobs=-1)]: Done 10 out of 10 | elapsed: 0.4s finished

Summary: Dataset without feature engineering resulted in  $R^2$  of 0.82 for Bagging Regressor.  
#9.VOTING - REGRESSOR

```
In [0]: from sklearn.ensemble import VotingRegressor
        voting_reg=VotingRegressor([('Bagging',bag_reg),('random_forest',random_reg),('KNeighbors',KNeighborsRegressor),
                                    ('Decision_tree',decision_tree_reg),('SVR',SVR),('Linear_SVM',LinearSVC),
                                    ('Linear_regression',lin_reg),('SGD_Regression',lin_reg_SGD)])
```

```
In [0]: from sklearn.ensemble import VotingRegressor
        voting_reg=VotingRegressor([('Bagging',bag_reg),('random_forest',random_reg),('KNeighbors',KNeighborsRegressor),
                                    ('Decision_tree',decision_tree_reg),('SVR',SVR),('Linear_SVM',LinearSVC),
                                    ('Linear_regression',lin_reg),('SGD_Regression',lin_reg_SGD)])

        voting_reg.fit(X_train,y_train)
        y_pred_voting_reg=voting_reg.predict(X_train)
        lin_perform_metrics(y_train,y_pred_voting_reg)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The default value of 'n_estimators' will increase from 10 to 100 in version 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of 'kernel' will change from 'rbf' to 'poly' in version 0.22. To avoid this warning, please pass the 'kernel' argument explicitly.", FutureWarning)
```

```
R2_Square: 0.8519626355760865
Explained_Variance_Score: 0.8535176678822731
Mean_absolute_error: 48.1274729452754
Median_absolute_error: 32.78111217725031
Max_error: 357.36210173430027
MSE: 4871.674088659906
RMSE: 69.7973788093787
```

Summary: Dataset without feature engineering resulted in  $R^2$  of 0.852 for Gradient Boosting Regressor algorithm.

```
In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2
```

```
cross_val_score(voting_reg,X_train,y_train,cv=cv).mean()
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The default value of 'n_estimators' will increase from 10 to 100 in version 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of 'kernel' will change from 'rbf' to 'poly' in version 0.22. To avoid this warning, please pass the 'kernel' argument explicitly.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The default value of 'n_estimators' will increase from 10 to 100 in version 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of 'kernel' will change from 'rbf' to 'poly' in version 0.22. To avoid this warning, please pass the 'kernel' argument explicitly.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The default value of 'n_estimators' will increase from 10 to 100 in version 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of 'kernel' will change from 'rbf' to 'poly' in version 0.22. To avoid this warning, please pass the 'kernel' argument explicitly.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The default value of 'n_estimators' will increase from 10 to 100 in version 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of 'kernel' will change from 'rbf' to 'poly' in version 0.22. To avoid this warning, please pass the 'kernel' argument explicitly.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The default value of 'n_estimators' will increase from 10 to 100 in version 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of 'kernel' will change from 'rbf' to 'poly' in version 0.22. To avoid this warning, please pass the 'kernel' argument explicitly.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The default value of 'n_estimators' will increase from 10 to 100 in version 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of 'kernel' will change from 'rbf' to 'poly' in version 0.22. To avoid this warning, please pass the 'kernel' argument explicitly.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The default value of 'n_estimators' will increase from 10 to 100 in version 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of 'kernel' will change from 'rbf' to 'poly' in version 0.22. To avoid this warning, please pass the 'kernel' argument explicitly.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The default value of 'n_estimators' will increase from 10 to 100 in version 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of 'kernel' will change from 'rbf' to 'poly' in version 0.22. To avoid this warning, please pass the 'kernel' argument explicitly.", FutureWarning)
```

```

/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The default
    "10 in version 0.20 to 100 in 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default val
    "avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The defa
    "10 in version 0.20 to 100 in 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default val
    "avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The defa
    "10 in version 0.20 to 100 in 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default val
    "avoid this warning.", FutureWarning)

```

```

Out[0]: 0.7666171999325883

```

```

In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2
        cross_val_score(voting_reg,X_poly,y_train,cv=cv).mean()

```

```

/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The defa
    "10 in version 0.20 to 100 in 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default val
    "avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The defa
    "10 in version 0.20 to 100 in 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default val
    "avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The defa
    "10 in version 0.20 to 100 in 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default val
    "avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The defa
    "10 in version 0.20 to 100 in 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default val
    "avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The defa
    "10 in version 0.20 to 100 in 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default val
    "avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The defa
    "10 in version 0.20 to 100 in 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default val
    "avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The defa
    "10 in version 0.20 to 100 in 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default val
    "avoid this warning.", FutureWarning)

```

```

"10 in version 0.20 to 100 in 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will be 'auto' in version 0.22. To silence this warning, you can set gamma='scale' or gamma='auto'.
"avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The default value of n_estimators will be 100 in version 0.22. To silence this warning, you can set n_estimators=100 or n_estimators=None.
"10 in version 0.20 to 100 in 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will be 'auto' in version 0.22. To silence this warning, you can set gamma='scale' or gamma='auto'.
"avoid this warning.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The default value of n_estimators will be 100 in version 0.22. To silence this warning, you can set n_estimators=100 or n_estimators=None.
"10 in version 0.20 to 100 in 0.22.", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will be 'auto' in version 0.22. To silence this warning, you can set gamma='scale' or gamma='auto'.
"avoid this warning.", FutureWarning)

```

```
Out[0]: 0.8227991645784952
```

#### #10.REGULARIZATION - RIDGE, LASSO & ELASTICNET

```
In [0]: from sklearn.linear_model import ElasticNet,Ridge,Lasso
```

```
#RIDGE
```

```
In [0]: lin_reg_ridge=Ridge(random_state=42)
lin_reg_ridge.get_params()
```

```
Out[0]: {'alpha': 1.0,
        'copy_X': True,
        'fit_intercept': True,
        'max_iter': None,
        'normalize': False,
        'random_state': 42,
        'solver': 'auto',
        'tol': 0.001}
```

```
In [0]: lin_reg_ridge.fit(X_train,y_train)
y_pred_ridge=lin_reg_ridge.predict(X_train)
```

```
lin_perform_metrics(y_train,y_pred_ridge)
```

```

R2_Square: 0.6333153782627892
Explained_Variance_Score: 0.6333153802117362
Mean_absolute_error: 79.26602188100725
Median_absolute_error: 56.63909008326195
Max_error: 476.67373277227273
MSE: 12014.48316689951
RMSE: 109.61059787675418

```

```
In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2
```

```
cross_val_score(lin_reg_ridge,X_train,y_train,cv=cv).mean()
```

Out[0]: 0.6303102044462026

In [0]: # SGD model applying on Train Data with feature engineering

```
lin_reg_ridge.fit(X_poly,y_train)
y_pred_ridge_poly=lin_reg_ridge.predict(X_poly)
lin_perform_metrics(y_train,y_pred_ridge_poly)
```


R2\_Square: 0.8579237229831651  
Explained\_Variance\_Score: 0.8579237231086724  
Mean\_absolute\_error: 49.5543129691245  
Median\_absolute\_error: 35.752660990657986  
Max\_error: 415.8863422684266  
MSE: 4655.153059180753  
RMSE: 68.22868208591423

In [0]: # Evaluate Cross Val Score on Dataset with added features  
cross\_val\_score(lin\_reg\_ridge,X\_poly,y\_train,cv=cv).mean()

Out[0]: 0.8374402883161943

In [0]: # LEARNING CURVE TO EVALUATE PERFORMANCE BASED ON BLOCKS OF SAMPLES

```
size=[100,1000,5000,7500,10000]
plot_learning_curve_samples(lin_reg_ridge, 'Linear Ridge_additional_feature', X_poly, y,
                             n_jobs=-1, train_sizes=size)
save_results_to = ''
plt.savefig(save_results_to + 'BEST_SOLUTION_ridge.png', dpi = 300)
```





```
In [0]: save_results_to = ''
        plt.savefig(save_results_to + 'BEST_SOLUTION.png', dpi = 300)
```

```
#LASSO
```

```
In [0]: lin_reg_lasso=Lasso(random_state=42)
        lin_reg_lasso.get_params()
```

```
Out[0]: {'alpha': 1.0,
         'copy_X': True,
         'fit_intercept': True,
         'max_iter': 1000,
         'normalize': False,
         'positive': False,
         'precompute': False,
         'random_state': 42,
         'selection': 'cyclic',
         'tol': 0.0001,
         'warm_start': False}
```

```
In [0]: lin_reg_lasso.fit(X_train,y_train)
        y_pred_lasso=lin_reg_lasso.predict(X_train)

        lin_perform_metrics(y_train,y_pred_lasso)
```

```
R2_Square: 0.6095893977896276
Explained_Variance_Score: 0.6095893977896277
Mean_absolute_error: 82.01205613748988
Median_absolute_error: 59.040115741902014
Max_error: 497.23437196973333
MSE: 12847.791651302496
RMSE: 113.34809946047837
```

```
In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2
```

```
        cross_val_score(lin_reg_lasso,X_train,y_train,cv=cv).mean()
```

```
Out[0]: 0.6067986473313757
```

```
In [0]: # model applying on Train Data with feature engineering
```

```
        lin_reg_lasso.fit(X_poly,y_train)
        y_pred_lasso_poly=lin_reg_lasso.predict(X_poly)
        lin_perform_metrics(y_train,y_pred_lasso_poly)
```

R2\_Square: 0.7576431951791133  
Explained\_Variance\_Score: 0.7576431951791134  
Mean\_absolute\_error: 65.87483367855849  
Median\_absolute\_error: 49.45996682807845  
Max\_error: 391.4252494352576  
MSE: 7975.576779895686  
RMSE: 89.30608478651209

In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2

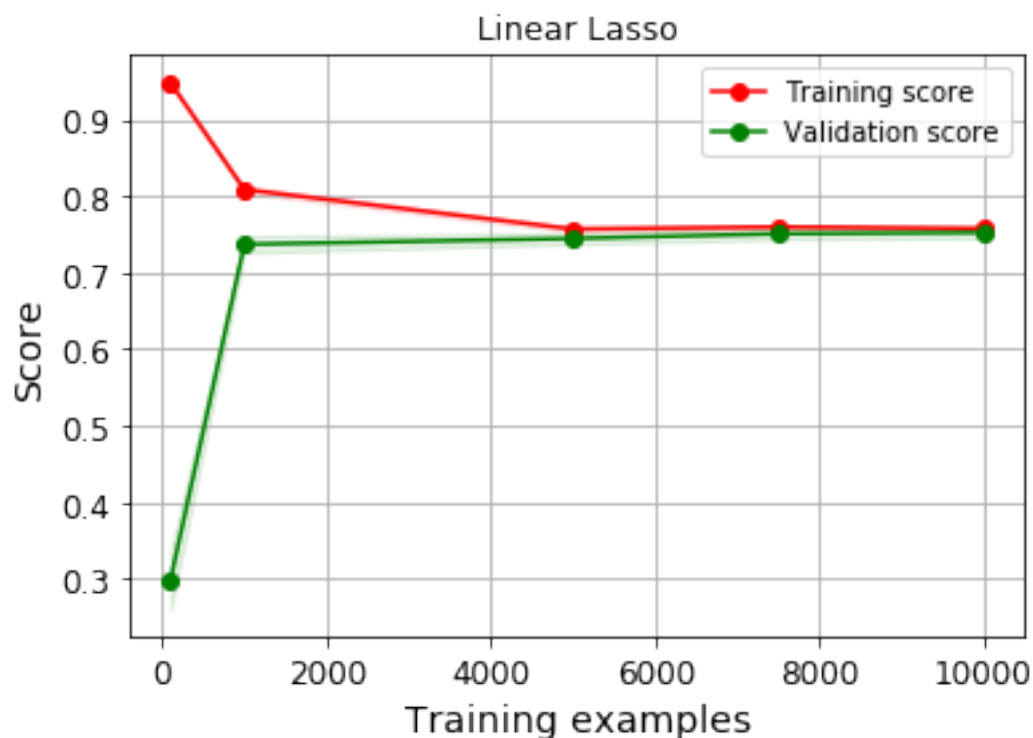
```
cross_val_score(lin_reg_lasso,X_poly,y_train,cv=cv).mean()
```

Out[0]: 0.7533756459859584

In [0]: # LEARNING CURVE TO EVALUATE PERFORMANCE BASED ON BLOCKS OF SAMPLES

```
size=[100,1000,5000,7500,10000]  
plot_learning_curve_samples(lin_reg_lasso, 'Linear Lasso', X_poly, y_train, ylim=None,  
                             n_jobs=-1, train_sizes=size)
```

Out[0]: <module 'matplotlib.pyplot' from '/usr/local/lib/python3.6/dist-packages/matplotlib/py'



#ELASTICNET

```
In [0]: lin_reg_elasticnet=ElasticNet(random_state=42)
lin_reg_elasticnet.get_params()
```

```
Out[0]: {'alpha': 1.0,
        'copy_X': True,
        'fit_intercept': True,
        'l1_ratio': 0.5,
        'max_iter': 1000,
        'normalize': False,
        'positive': False,
        'precompute': False,
        'random_state': 42,
        'selection': 'cyclic',
        'tol': 0.0001,
        'warm_start': False}
```

```
In [0]: lin_reg_elasticnet.fit(X_train,y_train)
y_pred_elasticnet=lin_reg_elasticnet.predict(X_train)

lin_perform_metrics(y_train,y_pred_elasticnet)
```

```
R2_Square: 0.29637353840005876
Explained_Variance_Score: 0.29637353840005876
Mean_absolute_error: 114.16234896593066
Median_absolute_error: 90.91706322131782
Max_error: 698.914785846821
MSE: 23155.22715776049
RMSE: 152.1684170837053
```

```
In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2
```

```
cross_val_score(lin_reg_elasticnet,X_train,y_train,cv=cv).mean()
```

```
Out[0]: 0.294726349109928
```

```
In [0]: # model applying on Train Data with feature engineering
```

```
lin_reg_elasticnet.fit(X_poly,y_train)
y_pred_elasticnet_poly=lin_reg_elasticnet.predict(X_poly)

lin_perform_metrics(y_train,y_pred_elasticnet_poly)
```

```
R2_Square: 0.46408325160365427
Explained_Variance_Score: 0.46408325160365427
Mean_absolute_error: 98.24432940799784
Median_absolute_error: 76.91663601426811
Max_error: 604.2315975364867
MSE: 17636.167375725083
RMSE: 132.80123258360626
```

```
In [0]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2
```

```
cross_val_score(lin_reg_elasticnet,X_poly,y_train,cv=cv).mean()
```


```
Out[0]: 0.46019225552513154
```

```
In [0]: # LEARNING CURVE TO EVALUATE PERFORMANCE BASED ON BLOCKS OF SAMPLES
```

```
size=[100,1000,5000,7500,10000]
```

```
plot_learning_curve_samples(lin_reg_elasticnet, 'Linear Elasticnet', X_poly, y_train, y  
                             n_jobs=-1, train_sizes=size)
```

```
Out[0]: <module 'matplotlib.pyplot' from '/usr/local/lib/python3.6/dist-packages/matplotlib/py
```



```
#BEST PARAMETERS
```

```
In [0]: # DETERMINE BEST PARAMETERS RESULTED ASSOCIATED WITH X_POLY DATA
```

```
alpha=[0.0001,0.001,0.01,0.1,1,10,100]
```

```
#estimators=[lin_reg_ridge,lin_reg_lasso,lin_reg_elasticnet]
```

```
param={'alpha':alpha}
```

```
best_parameters_poly(lin_reg_ridge,param,cv=cv)
```

Fitting 10 folds for each of 7 candidates, totalling 70 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.  
[Parallel(n_jobs=-1)]: Done 46 tasks      | elapsed: 1.1min  
[Parallel(n_jobs=-1)]: Done 70 out of 70 | elapsed: 1.4min finished
```

Best\_results: 0.8398230385631523

Best\_parameters: {'alpha': 10}

Best\_estimator: Ridge(alpha=10, copy\_X=True, fit\_intercept=True, max\_iter=None, normalize=False,  
random\_state=42, solver='auto', tol=0.001)

Summary: Dataset with feature engineering resulted in  $R^2$  of 0.8398 for Ridge linear algorithm.

#### #11.GRADIENT BOOSTING REGRESSOR

```
In [39]: from sklearn.ensemble import GradientBoostingRegressor  
lin_reg_boosting=GradientBoostingRegressor(max_depth=4,n_estimators=200)  
lin_reg_boosting.get_params()
```

```
Out[39]: {'alpha': 0.9,  
          'criterion': 'friedman_mse',  
          'init': None,  
          'learning_rate': 0.1,  
          'loss': 'ls',  
          'max_depth': 4,  
          'max_features': None,  
          'max_leaf_nodes': None,  
          'min_impurity_decrease': 0.0,  
          'min_impurity_split': None,  
          'min_samples_leaf': 1,  
          'min_samples_split': 2,  
          'min_weight_fraction_leaf': 0.0,  
          'n_estimators': 200,  
          'n_iter_no_change': None,  
          'presort': 'auto',  
          'random_state': None,  
          'subsample': 1.0,  
          'tol': 0.0001,  
          'validation_fraction': 0.1,  
          'verbose': 0,  
          'warm_start': False}
```

```
In [40]: lin_reg_boosting.fit(X_train,y_train)  
y_pred_boosting=lin_reg_boosting.predict(X_train)  
  
lin_perform_metrics(y_train,y_pred_boosting)
```

R2\_Square: 0.8462297206386022

Explained\_Variance\_Score: 0.8462297206386021

```
Mean_absolute_error: 49.761011736526136
Median_absolute_error: 33.143443142169346
Max_error: 417.6822383125031
MSE: 5038.309008445308
RMSE: 70.98104682551046
```

```
In [41]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2
```

```
cross_val_score(lin_reg_boosting,X_train,y_train,cv=cv).mean()
```


```
Out[41]: 0.8242665614621604
```

```
In [42]: lin_reg_boosting.fit(X_poly,y_train)
         y_pred_boosting_poly=lin_reg_boosting.predict(X_poly)

         lin_perform_metrics(y_train,y_pred_boosting_poly)
```

```
R2_Square: 0.8792380913903609
Explained_Variance_Score: 0.8792380913903609
Mean_absolute_error: 45.516768241866224
Median_absolute_error: 32.40517503925298
Max_error: 420.97269729588345
MSE: 3956.7842014191856
RMSE: 62.90297450374812
```

```
In [44]: plot_learning_curve_samples(lin_reg_boosting, 'GradientBoosting_additional_features',
                                     n_jobs=-1, train_sizes=size)
         save_results_to = ''
         plt.savefig(save_results_to + 'GradientBoosting', dpi = 300)
```



In [45]: # CROSS VALIDATION SCORE EVALUATION - MEAN OF R2

```
cross_val_score(lin_reg_boosting,X_poly,y_train,cv=cv).mean()
```

Out[45]: 0.8364402114863252

In [0]: # DETERMINE BEST PARAMETERS RESULTED ASSOCIATED WITH X\_TRAIN DATA

```
param ={'max_depth':[2,3,4],'n_estimators':[100,150,200]}
```

```
best_parameters_poly(lin_reg_boosting,param,cv=cv)
```

Fitting 10 folds for each of 9 candidates, totalling 90 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.

[Parallel(n\_jobs=-1)]: Done 46 tasks | elapsed: 9.7min

[Parallel(n\_jobs=-1)]: Done 90 out of 90 | elapsed: 26.7min finished

Best\_results: 0.8391519450108583

Best\_parameters: {'max\_depth': 4, 'n\_estimators': 200}

Best\_estimator: GradientBoostingRegressor(alpha=0.9, criterion='friedman\_mse', init=None, learning\_rate=0.1, loss='ls', max\_depth=4, max\_features=None, max\_leaf\_nodes=None, min\_impurity\_decrease=0.0, min\_impurity\_split=None,

```

min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=200,
n_iter_no_change=None, presort='auto',
random_state=None, subsample=1.0, tol=0.0001,
validation_fraction=0.1, verbose=0, warm_start=False)

```

```
In [0]: lin_perform_metrics(y_train,y_pred_boosting_poly)
```

Summary: Dataset with feature engineering resulted in  $R^2$  of 0.839 for Gradient Boosting Regressor algorithm.

## #12.NEURAL NETWORK

```
In [0]: def create_mlp(dim, regress=False):
        # define our MLP network
        model = Sequential()
        model.add(Dense(30, input_dim=dim, activation="relu"))
        model.add(Dense(10, activation="relu"))

        # check to see if the regression node should be added
        if regress:
            model.add(Dense(1, activation="linear"))

        # return our model
        return model

```

```
In [0]: from keras import backend as K
```

```

def coeff_determination(y_true, y_pred):
    SS_res = K.sum(K.square( y_true-y_pred ))
    SS_tot = K.sum(K.square( y_true - K.mean(y_true) ) )
    return ( 1 - SS_res/(SS_tot + K.epsilon()) )

```

```
In [0]: X_train_nn, X_valid_nn, y_train_nn, y_valid_nn = train_test_split(X_train, y_train, ran
```

```
In [0]: X_train_nn_p, X_valid_nn_p, y_train_nn_p, y_valid_nn_p = train_test_split(X_poly, y_tra
```

```
In [0]: from keras.optimizers import Adam,SGD
        from keras.models import Sequential
        from keras.layers.core import Dense

        model = create_mlp(X_poly.shape[1], regress=True)
        opt =Adam()
        model.compile(loss="mean_absolute_percentage_error", optimizer=opt,metrics=[coeff_dete
        history_1 = model.fit(X_train_nn_p, y_train_nn_p, epochs=30, validation_data=(X_valid_nn_p, y_valid_nn_p))

```

Train on 10427 samples, validate on 3476 samples

Epoch 1/30

10427/10427 [=====] - 2s 233us/step - loss: 86.0301 - coeff\_determina



Epoch 2/30  
10427/10427 [=====] - 1s 140us/step - loss: 64.1208 - coeff\_determina  
Epoch 3/30  
10427/10427 [=====] - 1s 142us/step - loss: 45.7853 - coeff\_determina  
Epoch 4/30  
10427/10427 [=====] - 2s 144us/step - loss: 38.1574 - coeff\_determina  
Epoch 5/30  
10427/10427 [=====] - 2s 144us/step - loss: 34.3664 - coeff\_determina  
Epoch 6/30  
10427/10427 [=====] - 1s 143us/step - loss: 32.1465 - coeff\_determina  
Epoch 7/30  
10427/10427 [=====] - 1s 143us/step - loss: 30.7885 - coeff\_determina  
Epoch 8/30  
10427/10427 [=====] - 1s 140us/step - loss: 29.7312 - coeff\_determina  
Epoch 9/30  
10427/10427 [=====] - 1s 141us/step - loss: 28.8263 - coeff\_determina  
Epoch 10/30  
10427/10427 [=====] - 1s 142us/step - loss: 28.0905 - coeff\_determina  
Epoch 11/30  
10427/10427 [=====] - 2s 150us/step - loss: 27.4140 - coeff\_determina  
Epoch 12/30  
10427/10427 [=====] - 2s 149us/step - loss: 26.8664 - coeff\_determina  
Epoch 13/30  
10427/10427 [=====] - 1s 143us/step - loss: 26.3851 - coeff\_determina  
Epoch 14/30  
10427/10427 [=====] - 2s 148us/step - loss: 26.0167 - coeff\_determina  
Epoch 15/30  
10427/10427 [=====] - 1s 143us/step - loss: 25.5707 - coeff\_determina  
Epoch 16/30  
10427/10427 [=====] - 2s 144us/step - loss: 25.2729 - coeff\_determina  
Epoch 17/30  
10427/10427 [=====] - 1s 143us/step - loss: 25.0296 - coeff\_determina  
Epoch 18/30  
10427/10427 [=====] - 1s 142us/step - loss: 24.6866 - coeff\_determina  
Epoch 19/30  
10427/10427 [=====] - 1s 144us/step - loss: 24.4063 - coeff\_determina  
Epoch 20/30  
10427/10427 [=====] - 1s 141us/step - loss: 24.0607 - coeff\_determina  
Epoch 21/30  
10427/10427 [=====] - 1s 141us/step - loss: 23.9037 - coeff\_determina  
Epoch 22/30  
10427/10427 [=====] - 1s 139us/step - loss: 23.5731 - coeff\_determina  
Epoch 23/30  
10427/10427 [=====] - 2s 149us/step - loss: 23.3791 - coeff\_determina  
Epoch 24/30  
10427/10427 [=====] - 2s 145us/step - loss: 23.2554 - coeff\_determina  
Epoch 25/30  
10427/10427 [=====] - 1s 141us/step - loss: 22.8347 - coeff\_determina

```

Epoch 26/30
10427/10427 [=====] - 1s 136us/step - loss: 22.8553 - coeff_determina
Epoch 27/30
10427/10427 [=====] - 1s 134us/step - loss: 22.4398 - coeff_determina
Epoch 28/30
10427/10427 [=====] - 1s 142us/step - loss: 22.2464 - coeff_determina
Epoch 29/30
10427/10427 [=====] - 1s 144us/step - loss: 22.1722 - coeff_determina
Epoch 30/30
10427/10427 [=====] - 2s 150us/step - loss: 21.8355 - coeff_determina

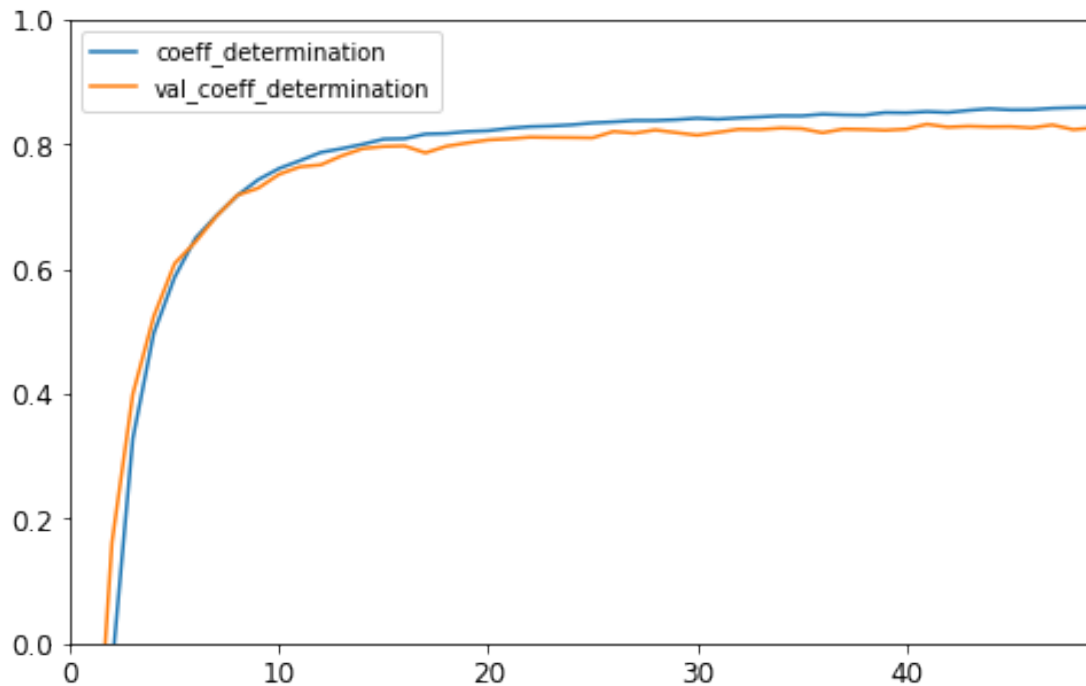
```

```

In [0]: df_1=pd.DataFrame(history_1.history)
        df_1[['coeff_determination','val_coeff_determination']].plot(figsize=(8,5))
        plt.gca().set_ylim(0, 1)

```

```
Out[0]: (0, 1)
```



```

In [0]: df_1[['coeff_determination','val_coeff_determination']].mean()

```

```

Out[0]: coeff_determination      0.715489
        val_coeff_determination  0.713778
        dtype: float64

```

```

In [0]: from tensorflow import keras

```

```

X_train_nn, X_valid_nn, y_train_nn, y_valid_nn = train_test_split(X_train, y_train, ran

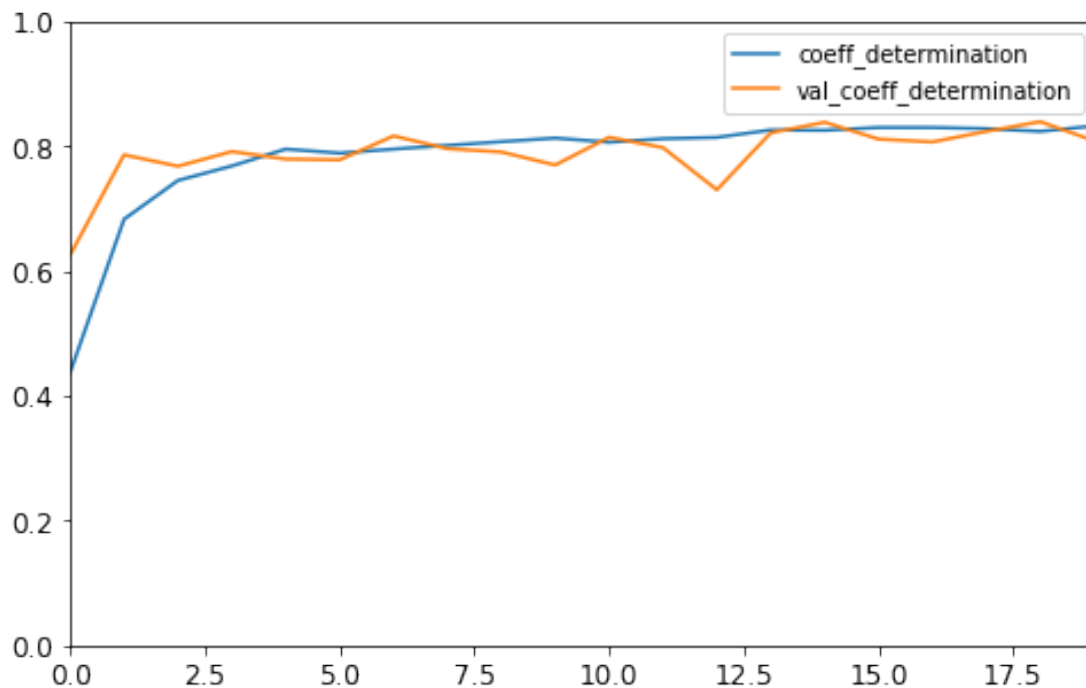
model = keras.models.Sequential([keras.layers.Dense(30, activation="relu", input_shape=
                                keras.layers.Dense(1))])
model.compile(loss='mean_squared_error', optimizer=keras.optimizers.SGD(lr=1e-3), metri
history = model.fit(X_train_nn, y_train_nn, epochs=20, validation_data=(X_valid_nn, y_v

Train on 10427 samples, validate on 3476 samples
Epoch 1/20
10427/10427 [=====] - 1s 103us/sample - loss: 16872.0308 - coeff_deter
Epoch 2/20
10427/10427 [=====] - 1s 87us/sample - loss: 9446.3624 - coeff_determ
Epoch 3/20
10427/10427 [=====] - 1s 88us/sample - loss: 7579.1897 - coeff_determ
Epoch 4/20
10427/10427 [=====] - 1s 85us/sample - loss: 6752.1117 - coeff_determ
Epoch 5/20
10427/10427 [=====] - 1s 90us/sample - loss: 6187.8912 - coeff_determ
Epoch 6/20
10427/10427 [=====] - 1s 87us/sample - loss: 6281.3251 - coeff_determ
Epoch 7/20
10427/10427 [=====] - 1s 88us/sample - loss: 6045.5160 - coeff_determ
Epoch 8/20
10427/10427 [=====] - 1s 87us/sample - loss: 5875.9883 - coeff_determ
Epoch 9/20
10427/10427 [=====] - 1s 87us/sample - loss: 5646.1375 - coeff_determ
Epoch 10/20
10427/10427 [=====] - 1s 88us/sample - loss: 5538.3928 - coeff_determ
Epoch 11/20
10427/10427 [=====] - 1s 88us/sample - loss: 5743.3635 - coeff_determ
Epoch 12/20
10427/10427 [=====] - 1s 88us/sample - loss: 5502.3247 - coeff_determ
Epoch 13/20
10427/10427 [=====] - 1s 88us/sample - loss: 5456.9718 - coeff_determ
Epoch 14/20
10427/10427 [=====] - 1s 87us/sample - loss: 5196.8318 - coeff_determ
Epoch 15/20
10427/10427 [=====] - 1s 88us/sample - loss: 5144.3832 - coeff_determ
Epoch 16/20
10427/10427 [=====] - 1s 88us/sample - loss: 5060.5831 - coeff_determ
Epoch 17/20
10427/10427 [=====] - 1s 87us/sample - loss: 5061.2031 - coeff_determ
Epoch 18/20
10427/10427 [=====] - 1s 87us/sample - loss: 5065.2998 - coeff_determ
Epoch 19/20
10427/10427 [=====] - 1s 89us/sample - loss: 5197.1768 - coeff_determ
Epoch 20/20
10427/10427 [=====] - 1s 89us/sample - loss: 4989.5875 - coeff_determ

```

```
In [0]: df=pd.DataFrame(history.history)
        df[['coeff_determination','val_coeff_determination']].plot(figsize=(8,5))
        plt.gca().set_ylim(0, 1)
```

```
Out[0]: (0, 1)
```



```
In [0]: df[['coeff_determination','val_coeff_determination']].mean()
```

```
Out[0]: coeff_determination      0.783398
        val_coeff_determination  0.789922
        dtype: float64
```

```
In [0]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20 entries, 0 to 19
Data columns (total 4 columns):
loss                20 non-null float64
coeff_determination 20 non-null float64
val_loss            20 non-null float64
val_coeff_determination 20 non-null float64
dtypes: float64(4)
memory usage: 720.0 bytes
```

```
In [0]: X_train_nn_poly, X_valid_nn_poly, y_train_nn_poly, y_valid_nn_poly = train_test_split(X_train_poly, y_train_poly,
                                                    test_size=0.2, random_state=0)

model = keras.models.Sequential([keras.layers.Dense(30, activation="relu", input_shape=X_train_poly.shape[1:]),
                                keras.layers.Dense(30, activation="relu"),
                                keras.layers.Dense(30, activation="relu"),
                                keras.layers.Dense(10, activation="softmax")])
model.compile(loss='mean_squared_error', optimizer=keras.optimizers.SGD(lr=1e-3), metrics=['accuracy'])
history = model.fit(X_train_nn_poly, y_train_nn_poly, epochs=20, validation_data=(X_valid_nn_poly, y_valid_nn_poly))
```

```
Epoch 1/20
10427/10427 [=====] - 1s 113us/sample - loss: 34096.7476 - coeff_determ
Epoch 2/20
10427/10427 [=====] - 1s 92us/sample - loss: 44840.6059 - coeff_determ
Epoch 3/20
10427/10427 [=====] - 1s 95us/sample - loss: 30905.2276 - coeff_determ
Epoch 4/20
10427/10427 [=====] - 1s 95us/sample - loss: 23653.2585 - coeff_determ
Epoch 5/20
10427/10427 [=====] - 1s 97us/sample - loss: 19396.2298 - coeff_determ
Epoch 6/20
10427/10427 [=====] - 1s 100us/sample - loss: 16897.0140 - coeff_determ
Epoch 7/20
10427/10427 [=====] - 1s 99us/sample - loss: 15755.7031 - coeff_determ
Epoch 8/20
10427/10427 [=====] - 1s 98us/sample - loss: 15676.6181 - coeff_determ
Epoch 9/20
10427/10427 [=====] - 1s 99us/sample - loss: 17344.6461 - coeff_determ
Epoch 10/20
10427/10427 [=====] - 1s 98us/sample - loss: 15226.0032 - coeff_determ
Epoch 11/20
10427/10427 [=====] - 1s 98us/sample - loss: 14184.5030 - coeff_determ
Epoch 12/20
10427/10427 [=====] - 1s 100us/sample - loss: 13432.2271 - coeff_determ
Epoch 13/20
10427/10427 [=====] - 1s 98us/sample - loss: 13243.3063 - coeff_determ
Epoch 14/20
10427/10427 [=====] - 1s 98us/sample - loss: 13049.4854 - coeff_determ
Epoch 15/20
10427/10427 [=====] - 1s 98us/sample - loss: 12767.7984 - coeff_determ
Epoch 16/20
10427/10427 [=====] - 1s 100us/sample - loss: 11629.0584 - coeff_determ
Epoch 17/20
10427/10427 [=====] - 1s 97us/sample - loss: 11143.6181 - coeff_determ
Epoch 18/20
10427/10427 [=====] - 1s 97us/sample - loss: 10398.5255 - coeff_determ
Epoch 19/20
10427/10427 [=====] - 1s 98us/sample - loss: 10107.0452 - coeff_determ
Epoch 20/20
10427/10427 [=====] - 1s 99us/sample - loss: 10713.9335 - coeff_determ
```

## #12.SUMMARY OF BEST RESULTS

1. SGD REGRESSOR: Best\_results with 2 degree polynomial feature : 0.8383639259139997/Best\_parameters: {'alpha': 0.001, 'penalty': 'elasticnet'}/ No overfitting
2. LINEAR REGRESSION : Best\_results with 2 degree polynomial feature : 0.8369484969213683/Best\_parameters: default value/No overfitting
3. LINEAR - SVR : Best\_results with 2 degree polynomial feature: 0.8067299323898737/Best\_parameters: {'C': 1, 'epsilon': 0.06}/No overfitting
4. SVR: Best\_results without feature engineering: 0.5666537147982439/ Best\_parameters: {'kernel': 'linear'}/ No overfitting
5. DECISION TREE: Best\_results with 2 degree polynomial feature : 0.7254413126025471/Best\_parameters: {'max\_depth': 15, 'max\_features': 0.71, 'min\_samples\_leaf': 0.001}/OVERFITTING
6. KNEIGHBORS: Best\_results without feature engineering: 0.6133113737857521/Best\_parameters: {'leaf\_size': 30, 'n\_neighbors': 5}/OVERFITTING
- 7.RANDOMFOREST: Best\_results without feature engineering: 0.8262559668916158/Best\_parameters: {'max\_features': 'auto', 'n\_estimators': 300}/OVERFITTING
8. BAGGING: Best\_results without feature engineering : 0.8215451437411845/Best\_parameters: default value/OVERFITTING
9. VOTING: Best\_results without feature engineering and default cv: 0.7666171999325883/OVERFITTING
10. REGULARIZATION - RIDGE : Best\_results without feature engineering: 0.8398230385631523/Best\_parameters: {'alpha': 10}/No overfitting
11. BOOSTING : Best\_results: 0.8391519450108583 / Best\_parameters: {'max\_depth': 4, 'n\_estimators': 200}/No overfitting
12. NEURAL NETWORK

CONCLUSION: SOME OF THE BEST MODELS ARE BOOSTING, RIDGE, SGD REGRESSOR, LINEAR REGRESSION - WITH CLOSE TO 0.84 R2 SCORE.

```
In [0]: print('BOOSTING : Best_results:' 0.8391519450108583 / Best_parameters: {'max_depth': 4
```

## #13.TEST MODELS #LINEAR REGRESSION

```
In [0]: y_pred_test_lr=lin_reg.predict(X_poly_test)
lin_perform_metrics(y_test,y_pred_test_lr)
```

```
R2_Square: 0.8347872765256323
Explained_Variance_Score: 0.8347902554913416
Mean_absolute_error: 53.58571465680695
Median_absolute_error: 37.83315012525807
Max_error: 374.6470617249621
MSE: 5524.231472061012
RMSE: 74.32517387844452
```

## 11 LINEAR SGD

```
In [0]: y_pred_test_SGD=lin_reg_SGD.predict(X_poly_test)
        lin_perform_metrics(y_test,y_pred_test_SGD)
```

```
R2_Square: 0.8371448664176161
Explained_Variance_Score: 0.8371509656517235
Mean_absolute_error: 53.493466003968535
Median_absolute_error: 37.747131123135084
Max_error: 365.60281341259076
MSE: 5445.400544238859
RMSE: 73.79295728075179
```

## 12 RIDGE

```
In [0]: y_pred_test_ridge=lin_reg_ridge.predict(X_poly_test)
        lin_perform_metrics(y_test,y_pred_test_ridge)
```

```
R2_Square: 0.8370721231284187
Explained_Variance_Score: 0.8370776260034953
Mean_absolute_error: 53.163349494977396
Median_absolute_error: 37.79986549276791
Max_error: 369.0919164720194
MSE: 5447.83286760424
RMSE: 73.80943616912569
```

### #GRADIENT BOOSTING

```
In [43]: y_pred_lin_reg_boosting=lin_reg_boosting.predict(X_poly_test)
         lin_perform_metrics(y_test,y_pred_lin_reg_boosting)
```

```
R2_Square: 0.8399968586862423
Explained_Variance_Score: 0.8400162555586852
Mean_absolute_error: 51.51212676840246
Median_absolute_error: 34.84228204869805
Max_error: 412.03051889537295
MSE: 5350.03824333917
```

RMSE: 73.14395561725637

#RANDOM FOREST

```
In [0]: y_pred_test_random_reg=random_reg.predict(X_poly_test)
        lin_perform_metrics(y_test,y_pred_test_random_reg)
```

R2\_Square: 0.8081877739351953  
Explained\_Variance\_Score: 0.8081936797109239  
Mean\_absolute\_error: 51.20738156885309  
Median\_absolute\_error: 29.099999999999998  
Max\_error: 480.9  
MSE: 6413.641235795454  
RMSE: 80.08521234157686

Conclusion: Select Linear Models for implementation  
Authors: Pranav, Hitendar, Suganthan,Mubasher & Asif